

Decision Tree Regression implementation of household electricity consumption dataset

Activities

- Data cleaning
- EDA and FE
- Model Building
- Accuracy check

Dataset:

<https://archive.ics.uci.edu/ml/datasets/individual+household+electric+power+consumption>

```
In [44]: import pandas as pd
import numpy as np
### Visualisation libraries
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
### For Q-Q Plot
import scipy.stats as stats
### To ignore warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [18]: df=pd.read_csv('household_power_consumption.txt', sep=';')
df.head()
```

C:\Users\subho\AppData\Local\Temp\ipykernel_8260\4214271158.py:1: DtypeWarning: Columns (2,3,4,5,6,7) have mixed types. Specify dtype option on import or set low_memory=False.
df=pd.read_csv('household_power_consumption.txt', sep=';')

```
Out[18]:
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1
0	16/12/2006	17:24:00	4.216	0.418	234.840	18.400	0.000
1	16/12/2006	17:25:00	5.360	0.436	233.630	23.000	0.000
2	16/12/2006	17:26:00	5.374	0.498	233.290	23.000	0.000
3	16/12/2006	17:27:00	5.388	0.502	233.740	23.000	0.000
4	16/12/2006	17:28:00	3.666	0.528	235.680	15.800	0.000

Basic Analysis

```
In [19]: df.shape
```

```
Out[19]: (2075259, 9)
```

```
In [20]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075259 entries, 0 to 2075258
Data columns (total 9 columns):
#   Column                                Dtype
---  -
0   Date                                  object
1   Time                                  object
2   Global_active_power                  object
3   Global_reactive_power                object
4   Voltage                              object
5   Global_intensity                     object
6   Sub_metering_1                       object
7   Sub_metering_2                       object
8   Sub_metering_3                       float64
dtypes: float64(1), object(8)
memory usage: 142.5+ MB
```

```
In [21]: ##### Taking 50000 samples from this huge dataset
data=df.sample(n=50000,replace=False)
data.reset_index(inplace=True)
data.head()
```

```
Out[21]:
```

	index	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_m
0	1565696	8/12/2009	00:20:00	0.272	0.000	244.340	1.200	
1	1941687	26/8/2010	02:51:00	0.342	0.208	239.760	1.600	
2	1579267	17/12/2009	10:31:00	1.440	0.106	243.860	5.800	
3	1652380	6/2/2010	05:04:00	0.250	0.000	244.300	1.000	
4	38696	12/1/2007	14:20:00	0.288	0.156	243.650	1.400	

```
In [23]: data.drop('index',axis=1,inplace=True)
data.shape
```

```
Out[23]: (50000, 9)
```

```
In [24]: ##### All the features are of object type except 'Sub_metering_3'. Need to convert all of
data['Date']=pd.to_datetime(data['Date'], format="%d/%m/%Y")
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Date                                  50000 non-null  datetime64[ns]
1   Time                                  50000 non-null  object
2   Global_active_power                  50000 non-null  object
3   Global_reactive_power                50000 non-null  object
4   Voltage                              50000 non-null  object
5   Global_intensity                     50000 non-null  object
6   Sub_metering_1                       50000 non-null  object
7   Sub_metering_2                       50000 non-null  object
8   Sub_metering_3                       49362 non-null  float64
dtypes: datetime64[ns](1), float64(1), object(7)
memory usage: 3.4+ MB
```

```
In [25]: ##### Creating day, month column to reduce dependancy on the 'Date' column
data['day']=data['Date'].dt.day
data['month']=data['Date'].dt.month
data.head()
```

Out [25]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub
0	2009-12-08	00:20:00	0.272	0.000	244.340	1.200	0.000	
1	2010-08-26	02:51:00	0.342	0.208	239.760	1.600	0.000	
2	2009-12-17	10:31:00	1.440	0.106	243.860	5.800	0.000	
3	2010-02-06	05:04:00	0.250	0.000	244.300	1.000	0.000	
4	2007-01-12	14:20:00	0.288	0.156	243.650	1.400	0.000	

In [26]: data['year']=data['Date'].dt.year
data.head()

Out [26]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub
0	2009-12-08	00:20:00	0.272	0.000	244.340	1.200	0.000	
1	2010-08-26	02:51:00	0.342	0.208	239.760	1.600	0.000	
2	2009-12-17	10:31:00	1.440	0.106	243.860	5.800	0.000	
3	2010-02-06	05:04:00	0.250	0.000	244.300	1.000	0.000	
4	2007-01-12	14:20:00	0.288	0.156	243.650	1.400	0.000	

In [27]: ##### dropping the date and time column
data.drop(['Date','Time'],axis=1,inplace=True)
data.head()

Out [27]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sul
0	0.272	0.000	244.340	1.200	0.000	0.000	
1	0.342	0.208	239.760	1.600	0.000	1.000	
2	1.440	0.106	243.860	5.800	0.000	0.000	
3	0.250	0.000	244.300	1.000	0.000	0.000	
4	0.288	0.156	243.650	1.400	0.000	2.000	

In [31]: ##### Checking null values in terms of '?'
for features in data.columns:
 c=0
 for values in data[features]:
 if values=='?':
 c+=1
 print(f'{features} has {c} null values')

```
Global_active_power has 638 null values
Global_reactive_power has 638 null values
Voltage has 638 null values
Global_intensity has 638 null values
Sub_metering_1 has 638 null values
Sub_metering_2 has 638 null values
Sub_metering_3 has 0 null values
day has 0 null values
month has 0 null values
year has 0 null values
```

```
In [37]: ##### handling null values
data.replace('?', np.nan, inplace=True)
data.isnull().sum()
```

```
Out[37]: Global_active_power      638
Global_reactive_power      638
Voltage                    638
Global_intensity           638
Sub_metering_1             638
Sub_metering_2             638
Sub_metering_3             638
day                        0
month                      0
year                       0
dtype: int64
```

```
In [38]: data.dropna(inplace=True)
data.isnull().sum()
```

```
Out[38]: Global_active_power      0
Global_reactive_power      0
Voltage                    0
Global_intensity           0
Sub_metering_1             0
Sub_metering_2             0
Sub_metering_3             0
day                        0
month                      0
year                       0
dtype: int64
```

```
In [39]: data.shape
```

```
Out[39]: (49362, 10)
```

```
In [40]: data.reset_index(inplace=True)
data.head()
```

```
Out[40]:
```

	index	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering
0	0	0.272	0.000	244.340	1.200	0.000	0.0
1	1	0.342	0.208	239.760	1.600	0.000	1.0
2	2	1.440	0.106	243.860	5.800	0.000	0.0
3	3	0.250	0.000	244.300	1.000	0.000	0.0
4	4	0.288	0.156	243.650	1.400	0.000	2.0

```
In [41]: data.drop('index', axis=1, inplace=True)
data.head()
```

Out[41]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sul
0	0.272	0.000	244.340	1.200	0.000	0.000	
1	0.342	0.208	239.760	1.600	0.000	1.000	
2	1.440	0.106	243.860	5.800	0.000	0.000	
3	0.250	0.000	244.300	1.000	0.000	0.000	
4	0.288	0.156	243.650	1.400	0.000	2.000	

In [42]:

```
#### we can now safely convert the features into numerical
for features in data.columns:
    data[features]=data[features].astype(float)
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 49362 entries, 0 to 49361
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Global_active_power    49362 non-null  float64
1   Global_reactive_power  49362 non-null  float64
2   Voltage                49362 non-null  float64
3   Global_intensity       49362 non-null  float64
4   Sub_metering_1         49362 non-null  float64
5   Sub_metering_2         49362 non-null  float64
6   Sub_metering_3         49362 non-null  float64
7   day                   49362 non-null  float64
8   month                 49362 non-null  float64
9   year                  49362 non-null  float64
dtypes: float64(10)
memory usage: 3.8 MB
```

In [43]:

```
data.describe().T
```

Out[43]:


	count	mean	std	min	25%	50%	75%	max
Global_active_power	49362.0	1.090352	1.062314	0.078	0.308	0.594	1.526	8.892
Global_reactive_power	49362.0	0.123314	0.112505	0.000	0.048	0.100	0.192	1.148
Voltage	49362.0	240.817051	3.268253	225.110	238.950	240.990	242.900	253.090
Global_intensity	49362.0	4.624898	4.468123	0.200	1.400	2.600	6.400	38.800
Sub_metering_1	49362.0	1.131255	6.179609	0.000	0.000	0.000	0.000	77.000
Sub_metering_2	49362.0	1.300089	5.813686	0.000	0.000	0.000	1.000	75.000
Sub_metering_3	49362.0	6.403245	8.419046	0.000	0.000	1.000	17.000	31.000
day	49362.0	15.711357	8.829163	1.000	8.000	16.000	23.000	31.000
month	49362.0	6.458024	3.423900	1.000	3.000	6.000	9.000	12.000
year	49362.0	2008.418318	1.123438	2006.000	2007.000	2008.000	2009.000	2010.000

Basic Graphical analysis


In [51]:

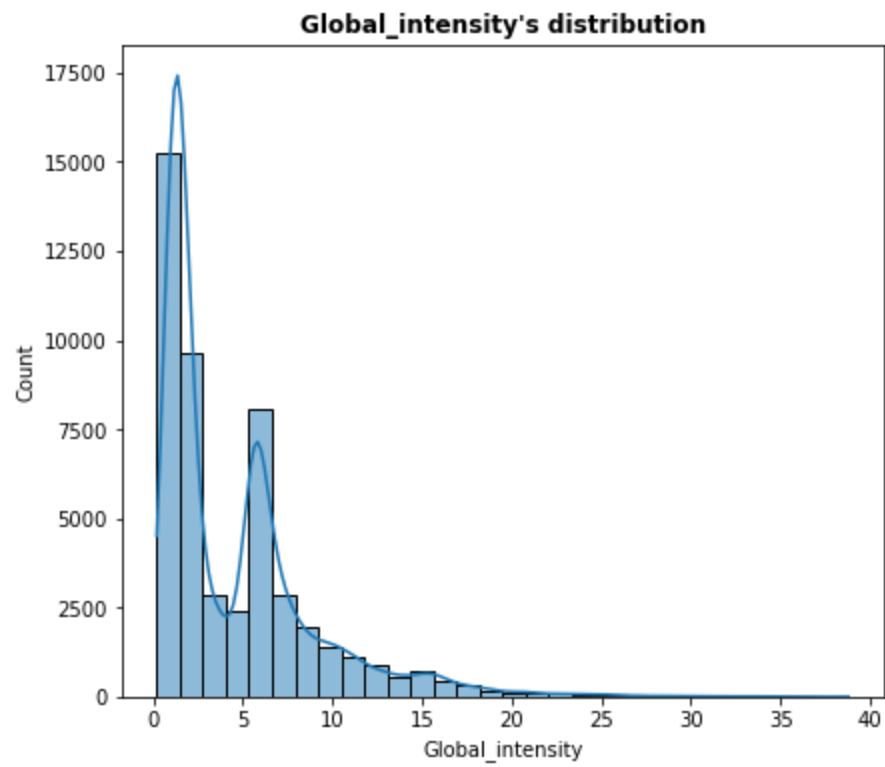

```
#### Checking the distribution of the dataset
for i in data:
    plt.figure(figsize=(15,6))
    plt.subplot(1,2,2)
    sns.histplot(data=data, x=i, kde=True, bins=30)
    plt.title("{}'s distribution".format(i),fontweight="bold")
```

Global_active_power's distribution




Global_reactive_power's distribution




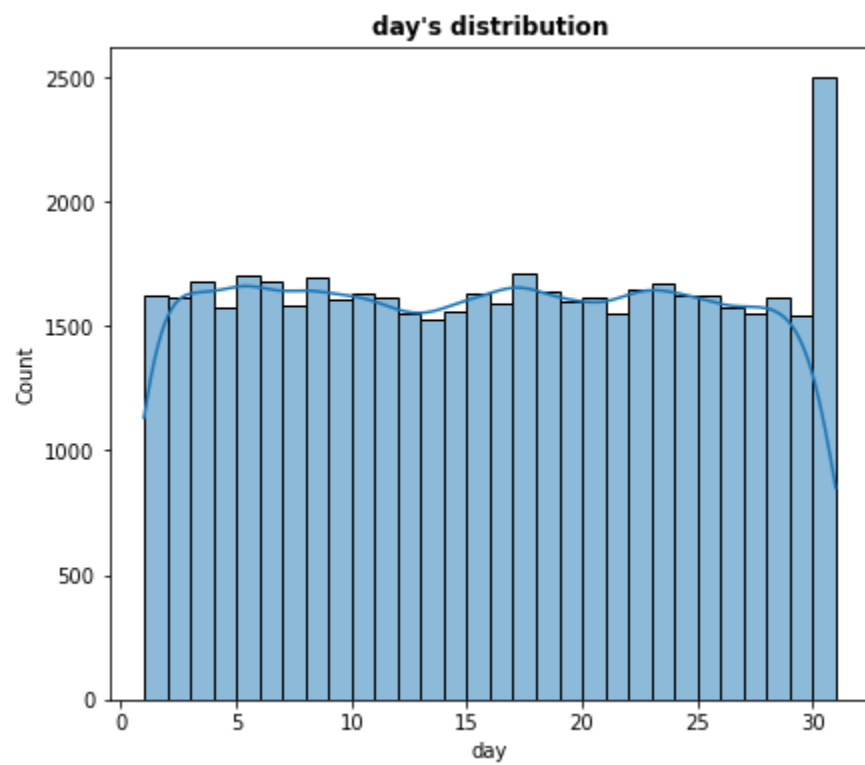
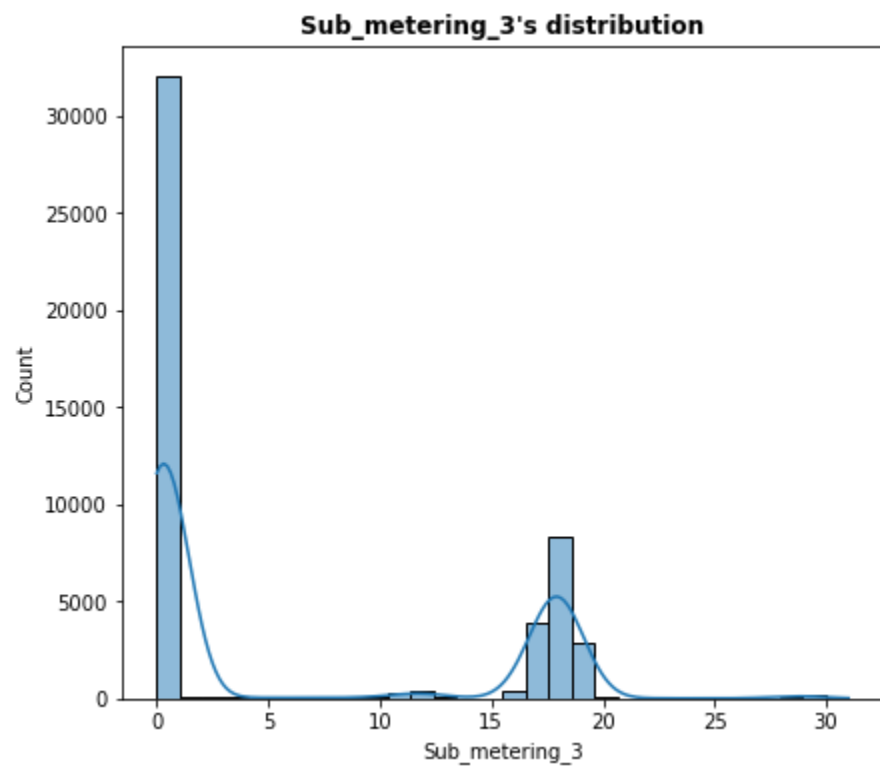




Sub_metering_1's distribution



Sub_metering_2's distribution











```
In [52]: ##### Crating a dependant feature 'Total_Power_Usage'
data['Total_Power_Usage']=data['Sub_metering_1']+data['Sub_metering_2']+data['Sub_metering_3']
data.head()
```



```
Out[52]:
```



	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	0.272	0.000	244.34	1.2	0.0	0.0	0.0
1	0.342	0.208	239.76	1.6	0.0	1.0	0.0
2	1.440	0.106	243.86	5.8	0.0	0.0	0.0
3	0.250	0.000	244.30	1.0	0.0	0.0	0.0
4	0.288	0.156	243.65	1.4	0.0	2.0	0.0



```
In [64]: ##### Comparing the features against dependant feature 'Total_Power_Usage'
for features in data:
    plot_num+=1
    if features!='Total_Power_Usage':
        plt.figure(figsize=(10,6))
        plt.scatter(x=data[features],y=data['Total_Power_Usage'],c='blue')
        plt.xlabel(features)
        plt.ylabel('Total_Power_Usage')
        plt.tight_layout()
        plot_num+=1
```












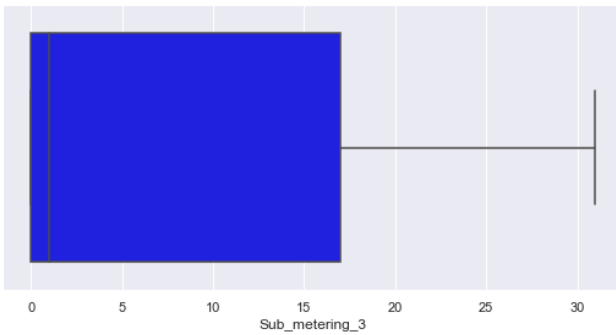
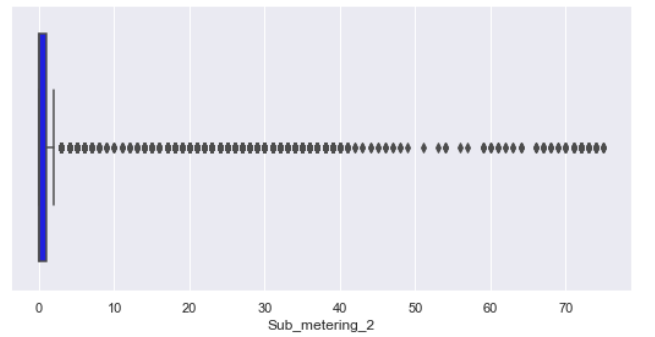
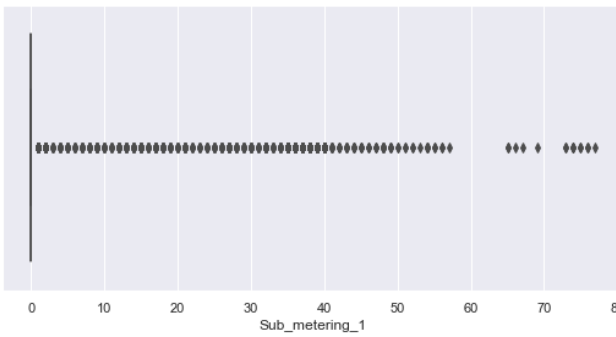

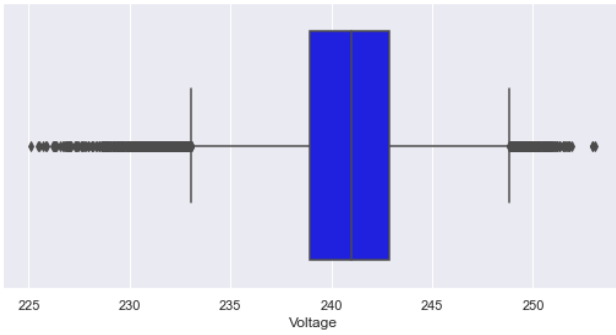


```
In [65]: ##### Checking correlation between the features
data.corr()
```

Out[65]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1
Global_active_power	1.000000	0.247642	-0.403876	0.998900	0.492156
Global_reactive_power	0.247642	1.000000	-0.114857	0.266809	0.134770
Voltage	-0.403876	-0.114857	1.000000	-0.415413	-0.196462
Global_intensity	0.998900	0.266809	-0.415413	1.000000	0.496834
Sub_metering_1	0.492156	0.134770	-0.196462	0.496834	1.000000
Sub_metering_2	0.438064	0.138916	-0.170027	0.444213	0.072243
Sub_metering_3	0.641381	0.085111	-0.277035	0.629476	0.106653
day	0.006372	0.000947	0.023519	0.005849	-0.002583
month	-0.031237	0.023582	-0.018293	-0.029994	-0.006605
year	-0.036686	0.040599	0.256456	-0.041184	-0.014463
Total_Power_Usage	0.848660	0.182044	-0.350258	0.845913	0.579077

Handling outliers

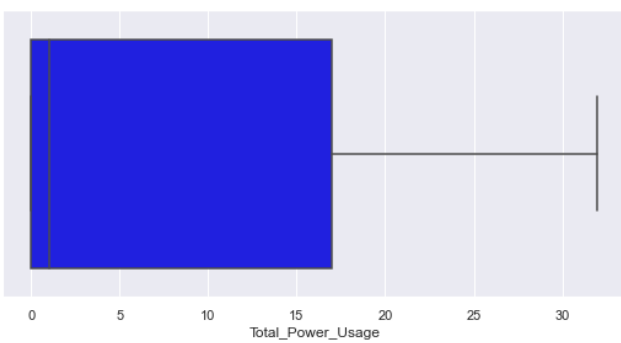







```
In [85]: ##### Visualising outliers using boxplot
plt.figure(figsize=(20,30))
for features in enumerate(data.columns):
    if features[1] not in ['day', 'month', 'year']:
        plt.subplot(6,2,features[0]+1)
        sns.set(rc={'figure.figsize':(10,8)})
        sns.boxplot(data=data,x=features[1],color='blue')
        plt.xlabel(features[1])
```

```
In [102... ##### From the above plot it is very evident that the data contains a lot of outliers. W
def outlier_handle(feature):
    q1=data[feature].quantile(0.25)
    q3=data[feature].quantile(0.75)
    IQR=q3-q1
    lower_limit=q1-1.5*IQR
    upper_limit=q3+1.5*IQR
    index_to_drop=list(data[(data[feature]<lower_limit)|(data[feature]>upper_limit)].index)
    data.drop(index_to_drop,inplace=True)
    return 'Outliers Removed'
```

```
In [105... ##### We are going to pass the features one by one to our function to remove outliers
for features in data.columns:
    if features not in ['day', 'month', 'year']:
        outlier_handle(features)
```

```
In [106... ##### Checking the boxplot post outlier removal
plt.figure(figsize=(20,30))
for features in enumerate(data.columns):
    if features[1] not in ['day', 'month', 'year']:
        plt.subplot(6,2,features[0]+1)
        sns.set(rc={'figure.figsize':(10,8)})
        sns.boxplot(data=data,x=features[1],color='blue')
        plt.xlabel(features[1])
```



```
In [107... ##### The plots look much better now. We can focus on model building now.
```

```
In [110... ##### Removing duplicate data
data.duplicated().sum()
```

```
Out[110]: 17
```

```
In [111... data.drop_duplicates(inplace=True)
data.reset_index(inplace=True)
data.duplicated().sum()
```

```
Out[111]: 0
```

```
In [113... data.drop('index',axis=1,inplace=True)
data.head()
```

```
Out[113]:
```

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	0.272	0.000	244.34	1.2	0.0	0.0	0.0
1	0.342	0.208	239.76	1.6	0.0	1.0	0.0
2	1.440	0.106	243.86	5.8	0.0	0.0	0.0
3	0.250	0.000	244.30	1.0	0.0	0.0	0.0
4	0.288	0.156	243.65	1.4	0.0	2.0	0.0

Model Building

```
In [114... ##### Splitting Dependent and Independent columns
X=data.drop('Total_Power_Usage',axis=1)
y=data['Total_Power_Usage']
X.head()
```

```
Out[114]:
```

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	0.272	0.000	244.34	1.2	0.0	0.0	0.0
1	0.342	0.208	239.76	1.6	0.0	1.0	0.0
2	1.440	0.106	243.86	5.8	0.0	0.0	0.0
3	0.250	0.000	244.30	1.0	0.0	0.0	0.0
4	0.288	0.156	243.65	1.4	0.0	2.0	0.0

```
In [115... y.head()
```

```
Out[115]:
```

0	1.0
1	2.0
2	19.0
3	1.0
4	2.0

Name: Total_Power_Usage, dtype: float64

```
In [116... ##### Splitting the training and test data
from sklearn.model_selection import train_test_split,GridSearchCV
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_state=10)
```

```
In [117... from sklearn.tree import DecisionTreeRegressor
model=DecisionTreeRegressor()
```

Out[118]: ▼ DecisionTreeRegressor
DecisionTreeRegressor()

In [120... model.score(X_train,y_train)

Out[120]: 1.0

In [121... **from** sklearn **import** tree
import matplotlib.pyplot **as** plt
fig=plt.figure(figsize=(25,15))
tree.plot_tree(model, max_depth=4, filled=True)

```

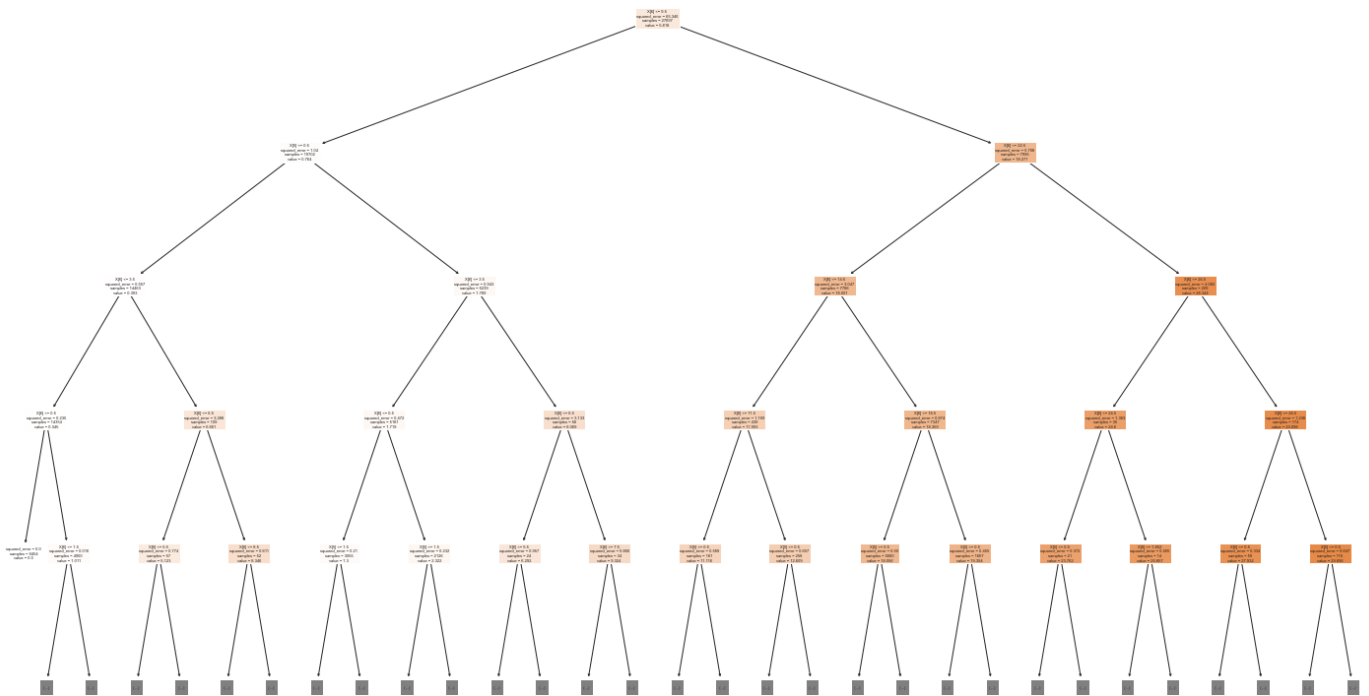
Out[121]: [Text(0.4774590163934426, 0.9166666666666666, 'X[6] <= 9.5\nsquared_error = 65.346\nsam
ples = 27697\nvalue = 5.818'),
Text(0.21721311475409835, 0.75, 'X[5] <= 0.5\nsquared_error = 1.04\nsamples = 19702\nv
alue = 0.764'),
Text(0.09016393442622951, 0.5833333333333334, 'X[6] <= 3.5\nsquared_error = 0.557\nsam
ples = 14463\nvalue = 0.393'),
Text(0.03278688524590164, 0.4166666666666667, 'X[6] <= 0.5\nsquared_error = 0.235\nsam
ples = 14354\nvalue = 0.345'),
Text(0.01639344262295082, 0.25, 'squared_error = 0.0\nsamples = 9454\nvalue = 0.0'),
Text(0.04918032786885246, 0.25, 'X[6] <= 1.5\nsquared_error = 0.016\nsamples = 4900\nv
alue = 1.011'),
Text(0.03278688524590164, 0.08333333333333333, '\n (...) \n'),
Text(0.06557377049180328, 0.08333333333333333, '\n (...) \n'),
Text(0.14754098360655737, 0.4166666666666667, 'X[6] <= 6.5\nsquared_error = 3.288\nsam
ples = 109\nvalue = 6.661'),
Text(0.11475409836065574, 0.25, 'X[6] <= 5.5\nsquared_error = 0.774\nsamples = 57\nval
ue = 5.123'),
Text(0.09836065573770492, 0.08333333333333333, '\n (...) \n'),
Text(0.13114754098360656, 0.08333333333333333, '\n (...) \n'),
Text(0.18032786885245902, 0.25, 'X[6] <= 8.5\nsquared_error = 0.611\nsamples = 52\nval
ue = 8.346'),
Text(0.16393442622950818, 0.08333333333333333, '\n (...) \n'),
Text(0.19672131147540983, 0.08333333333333333, '\n (...) \n'),
Text(0.3442622950819672, 0.5833333333333334, 'X[6] <= 3.5\nsquared_error = 0.943\nsamp
les = 5239\nvalue = 1.789'),
Text(0.2786885245901639, 0.4166666666666667, 'X[6] <= 0.5\nsquared_error = 0.472\nsamp
les = 5181\nvalue = 1.719'),
Text(0.2459016393442623, 0.25, 'X[5] <= 1.5\nsquared_error = 0.21\nsamples = 3055\nval
ue = 1.3'),
Text(0.22950819672131148, 0.08333333333333333, '\n (...) \n'),
Text(0.26229508196721313, 0.08333333333333333, '\n (...) \n'),
Text(0.3114754098360656, 0.25, 'X[5] <= 1.5\nsquared_error = 0.232\nsamples = 2126\nva
lue = 2.322'),
Text(0.29508196721311475, 0.08333333333333333, '\n (...) \n'),
Text(0.32786885245901637, 0.08333333333333333, '\n (...) \n'),
Text(0.4098360655737705, 0.4166666666666667, 'X[6] <= 6.5\nsquared_error = 3.133\nsamp
les = 58\nvalue = 8.069'),
Text(0.3770491803278688, 0.25, 'X[6] <= 5.5\nsquared_error = 0.957\nsamples = 24\nvalu
e = 6.292'),
Text(0.36065573770491804, 0.08333333333333333, '\n (...) \n'),
Text(0.39344262295081966, 0.08333333333333333, '\n (...) \n'),
Text(0.4426229508196721, 0.25, 'X[6] <= 7.5\nsquared_error = 0.866\nsamples = 34\nvalu
e = 9.324'),
Text(0.4262295081967213, 0.08333333333333333, '\n (...) \n'),
Text(0.45901639344262296, 0.08333333333333333, '\n (...) \n'),
Text(0.7377049180327869, 0.75, 'X[6] <= 22.5\nsquared_error = 5.798\nsamples = 7995\nv
alue = 18.271'),
Text(0.6065573770491803, 0.5833333333333334, 'X[6] <= 14.5\nsquared_error = 3.047\nsam
ples = 7786\nvalue = 18.001'),
Text(0.5409836065573771, 0.4166666666666667, 'X[6] <= 11.5\nsquared_error = 1.169\nsam
ples = 439\nvalue = 11.993'),
Text(0.5081967213114754, 0.25, 'X[5] <= 0.5\nsquared_error = 0.589\nsamples = 181\nval
ue = 11.116'),
Text(0.4918032786885246, 0.08333333333333333, '\n (...) \n'),
Text(0.5245901639344263, 0.08333333333333333, '\n (...) \n'),
Text(0.5737704918032787, 0.25, 'X[5] <= 0.5\nsquared_error = 0.657\nsamples = 258\nval
ue = 12.609'),
Text(0.5573770491803278, 0.08333333333333333, '\n (...) \n'),
Text(0.5901639344262295, 0.08333333333333333, '\n (...) \n'),
Text(0.6721311475409836, 0.4166666666666667, 'X[6] <= 18.5\nsquared_error = 0.874\nsam
ples = 7347\nvalue = 18.359'),
Text(0.639344262295082, 0.25, 'X[5] <= 0.5\nsquared_error = 0.59\nsamples = 5680\nvalu
e = 18.056'),
Text(0.6229508196721312, 0.08333333333333333, '\n (...) \n'),

```

```

Text(0.6557377049180327, 0.08333333333333333, '\n (...) \n'),
Text(0.7049180327868853, 0.25, 'X[5] <= 0.5\nsquared_error = 0.455\nsamples = 1667\nvalue = 19.394'),
Text(0.6885245901639344, 0.08333333333333333, '\n (...) \n'),
Text(0.7213114754098361, 0.08333333333333333, '\n (...) \n'),
Text(0.8688524590163934, 0.5833333333333334, 'X[6] <= 26.5\nsquared_error = 4.082\nsamples = 209\nvalue = 28.344'),
Text(0.8032786885245902, 0.4166666666666667, 'X[6] <= 24.5\nsquared_error = 1.383\nsamples = 35\nvalue = 24.6'),
Text(0.7704918032786885, 0.25, 'X[5] <= 0.5\nsquared_error = 0.372\nsamples = 21\nvalue = 23.762'),
Text(0.7540983606557377, 0.08333333333333333, '\n (...) \n'),
Text(0.7868852459016393, 0.08333333333333333, '\n (...) \n'),
Text(0.8360655737704918, 0.25, 'X[0] <= 1.852\nsquared_error = 0.265\nsamples = 14\nvalue = 25.857'),
Text(0.819672131147541, 0.08333333333333333, '\n (...) \n'),
Text(0.8524590163934426, 0.08333333333333333, '\n (...) \n'),
Text(0.9344262295081968, 0.4166666666666667, 'X[6] <= 28.5\nsquared_error = 1.238\nsamples = 174\nvalue = 29.098'),
Text(0.9016393442622951, 0.25, 'X[5] <= 0.5\nsquared_error = 0.334\nsamples = 59\nvalue = 27.932'),
Text(0.8852459016393442, 0.08333333333333333, '\n (...) \n'),
Text(0.9180327868852459, 0.08333333333333333, '\n (...) \n'),
Text(0.9672131147540983, 0.25, 'X[5] <= 0.5\nsquared_error = 0.647\nsamples = 115\nvalue = 29.696'),
Text(0.9508196721311475, 0.08333333333333333, '\n (...) \n'),
Text(0.9836065573770492, 0.08333333333333333, '\n (...) \n')]

```



```
In [122... y_predict=model.predict(X_test)
```

```
In [123... from sklearn.metrics import r2_score
```

```
In [128... r2=r2_score(y_test,y_predict)
print('Our model has', round(r2*100,3), '% accuracy')
```

Our model has 99.998 % accuracy

In []:

