

DTR_Household_Power_Consumption_Regression_Problem

November 18, 2022

```
# **
```

Regression Problem (Decision Trees Regressor and Extra Trees Regressor and Hyperparameter Tuning)

* To predict daily power consumption**

Datatable : To read large dataset

```
[1]: # !pip install datatable
```

Import required libraries

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
# import datatable as dt
%matplotlib inline
```

Complete dataset is available on my GitHub * GitHub Link:

https://github.com/subhashdixit/Regression_Model_Tasks/tree/main/Household_Power_Consumption_Regression_Problem

```
[3]: # # Datatable python library
# datatable_df = dt.fread("/content/drive/MyDrive/FSDS_Job_Guarantee/Tasks/
↳Regression_Problems/Household_Power_Consumption_Regression_Problem/
↳household_power_consumption.txt")
```

Convert Datatable into Pandas Dataframe

```
[4]: # df = datatable_df.to_pandas()
```

```
[5]: # df.head()
```

```
[6]: # df.columns
```

```
[7]: # df.shape
```

1 Sampling

- Take 100000 samples out of 2075259

```
[8]: # df=df.sample(100000).reset_index().drop('index',axis=1)
```

```
[9]: # df.head()
```

Store sample taken into csv for faster operation in future and also to avoid sampling every time. If we do sampling every time then our results will be impacted

```
[10]: # from google.colab import files
# df.to_csv('household_power_consumption_100000_samples.csv')
# files.download('household_power_consumption_100000_samples.csv')
```

Read Data From GitHub

Sample dataset is also available on my GitHub * GitHub Link:

https://raw.githubusercontent.com/subhashdixit/Regression_Model_Tasks/main/Household_Power_Consumption

```
[11]: url = 'https://raw.githubusercontent.com/subhashdixit/Regression_Model_Tasks/
↳main/Household_Power_Consumption_Regression_Problem/
↳household_power_consumption_100000_samples.csv'
df = pd.read_csv(url)
```

```
[12]: df.head()
```

```
[12]: Unnamed: 0      Date      Time Global_active_power Global_reactive_power \
0          0  23/9/2007  20:53:00           4.548           0.048
1          1  16/6/2008  05:44:00           0.332           0.266
2          2  20/7/2010  11:26:00           0.456           0.300
3          3  21/5/2008  07:58:00           1.402           0.110
4          4  16/6/2007  13:07:00           1.662           0.114
```

```
      Voltage Global_intensity Sub_metering_1 Sub_metering_2 Sub_metering_3
0  233.920           19.400           36.000           0.000           17.0
1  240.490           1.800           0.000           1.000           1.0
2  241.510           2.400           0.000           0.000           0.0
3  238.880           5.800           0.000           0.000          18.0
4  240.190           7.000           0.000           1.000          17.0
```

Data Set Information:

This archive contains 2075259 measurements gathered in a house located in Sceaux (7km of Paris, France) between December 2006 and November 2010 (47 months) * **We have taken 100000 samples only to predict power consumption**

Notes: 1. (global_active_power*1000/60 - sub_metering_1 - sub_metering_2 - sub_metering_3) represents the active energy consumed every minute (in watt hour) in the household by electrical equipment not measured in sub-meterings 1, 2 and 3 2. The dataset contains some missing values

in the measurements (nearly 1,25% of the rows). All calendar timestamps are present in the dataset but for some timestamps, the measurement values are missing: a missing value is represented by the absence of value between two consecutive semi-colon attribute separators. For instance, the dataset shows missing values on April 28, 2007

Attribute Information:

1. date:
 - Date in format dd/mm/yyyy
2. time:
 - time in format hh:mm:ss
3. global_active_power:
 - household global minute-averaged active power (in kilowatt)
4. global_reactive_power:
 - household global minute-averaged reactive power (in kilowatt)
5. voltage:
 - minute-averaged voltage (in volt)
6. global_intensity:
 - household global minute-averaged current intensity (in ampere)
7. sub_metering_1:
 - energy sub-metering No. 1 (in watt-hour of active energy). It corresponds to the kitchen, containing mainly a dishwasher, an oven and a microwave (hot plates are not electric but gas powered).
8. sub_metering_2:
 - energy sub-metering No. 2 (in watt-hour of active energy). It corresponds to the laundry room, containing a washing-machine, a tumble-drier, a refrigerator and a light.
9. sub_metering_3:
 - energy sub-metering No. 3 (in watt-hour of active energy). It corresponds to an electric water-heater and an air-conditioner.

```
[13]: df.columns
```

```
[13]: Index(['Unnamed: 0', 'Date', 'Time', 'Global_active_power',  
        'Global_reactive_power', 'Voltage', 'Global_intensity',  
        'Sub_metering_1', 'Sub_metering_2', 'Sub_metering_3'],  
        dtype='object')
```

Drop "Unnamed: 0" column because it is of no use

```
[14]: df.drop(['Unnamed: 0'], axis = 1, inplace = True)
```

2 EDA

Information about the dataset

```
[15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   Date                        100000 non-null object
1   Time                        100000 non-null object
2   Global_active_power        100000 non-null object
3   Global_reactive_power      100000 non-null object
4   Voltage                    100000 non-null object
5   Global_intensity           100000 non-null object
6   Sub_metering_1             100000 non-null object
7   Sub_metering_2             100000 non-null object
8   Sub_metering_3             98733 non-null  float64
dtypes: float64(1), object(8)
memory usage: 6.9+ MB
```

We will do our analysis on the basis of Daily Data and ignore time column

```
[16]: df['Date'] = pd.to_datetime(df['Date'])
```

```
[17]: df.drop(['Time'], axis = 1, inplace = True)
```

```
[18]: df.head()
```

```
[18]:
```

	Date	Global_active_power	Global_reactive_power	Voltage	\
0	2007-09-23	4.548	0.048	233.920	
1	2008-06-16	0.332	0.266	240.490	
2	2010-07-20	0.456	0.300	241.510	
3	2008-05-21	1.402	0.110	238.880	
4	2007-06-16	1.662	0.114	240.190	

	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	19.400	36.000	0.000	17.0
1	1.800	0.000	1.000	1.0
2	2.400	0.000	0.000	0.0
3	5.800	0.000	0.000	18.0
4	7.000	0.000	1.000	17.0

```
[19]: df.isnull().sum()
```

```
[19]: Date          0
      Global_active_power  0
      Global_reactive_power  0
      Voltage          0
      Global_intensity  0
      Sub_metering_1    0
      Sub_metering_2    0
      Sub_metering_3    1267
      dtype: int64
```

```
[20]: df.duplicated().sum()
```

```
[20]: 1326
```

Drop duplicates data

```
[21]: df.drop_duplicates(inplace = True)
```

```
[22]: df.columns
```

```
[22]: Index(['Date', 'Global_active_power', 'Global_reactive_power', 'Voltage',
          'Global_intensity', 'Sub_metering_1', 'Sub_metering_2',
          'Sub_metering_3'],
          dtype='object')
```

```
[23]: df['Sub_metering_1'].unique()
```

```
[23]: array(['36.000', '0.000', '2.000', '?', '3.000', '1.000', '37.000',
          '12.000', '13.000', '7.000', '33.000', '38.000', '39.000', '0.0',
          '40.000', '19.000', '31.000', '14.000', '29.000', '42.000',
          '35.000', '8.000', '5.000', '27.000', '18.000', '10.000', '16.000',
          '1.0', '9.000', '6.000', '21.000', '25.000', '43.000', '34.000',
          '52.000', '22.000', '11.000', '49.000', '41.000', '51.000',
          '24.000', '32.000', '44.000', '23.000', '47.000', '30.000',
          '15.000', '76.000', '53.000', '26.000', '20.000', '37.0', '28.000',
          '17.000', '2.0', '4.000', '45.000', '79.000', '38.0', '48.000',
          '78.000', '73.000', '75.000', '12.0', '74.000', '77.000', '46.000',
          '70.000', '54.000', '71.000', '67.000', '50.000', '60.000',
          '72.000', '9.0', '55.000'], dtype=object)
```

```
[24]: df.replace('?', np.nan, inplace=True)
```

```
[25]: df.isnull().sum()
```

```
[25]: Date          0
      Global_active_power  29
      Global_reactive_power  29
```

```
Voltage                29
Global_intensity       29
Sub_metering_1         29
Sub_metering_2         29
Sub_metering_3         29
dtype: int64
```

```
[26]: df.dropna(how = 'any', inplace = True)
```

```
[27]: # df.fillna(df.median().round(1), inplace=True)
```

```
[28]: df.isnull().sum()
```

```
[28]: Date                0
Global_active_power      0
Global_reactive_power    0
Voltage                  0
Global_intensity         0
Sub_metering_1           0
Sub_metering_2           0
Sub_metering_3           0
dtype: int64
```

Convert data to float datatype because all values are in decimal

```
[29]: convert_data = {'Global_active_power' : 'float64', 'Global_reactive_power' : 'float64', 'Voltage' : 'float64',
    ↪ 'Global_intensity' : 'float64', 'Sub_metering_1' : 'float64',
    ↪ 'Sub_metering_2' : 'float64',
    ↪ 'Sub_metering_3' : 'float64'}
df = df.astype(convert_data)
```

Take date wise data only

```
[30]: df= df.groupby('Date').sum()
```

```
[31]: df.reset_index(inplace = True)
```

```
[32]: df['year']=df['Date'].dt.year
df['month']=df['Date'].dt.month
```

```
[33]: df.groupby('year').sum()
```

```
[33]:      Global_active_power  Global_reactive_power  Voltage \
year
2006              2089.234              145.794    266246.20
2007              27953.588             2934.368   6026508.66
```

2008	27070.042	2976.726	6106483.96
2009	26946.336	3304.390	6046970.02
2010	23376.046	2867.062	5311922.51

	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3	month
year					
2006	8834.6	1319.0	2443.0	7970.0	192
2007	119247.4	32306.0	41834.0	144122.0	2378
2008	114927.2	27497.0	32388.0	152028.0	2384
2009	113853.8	27616.0	27921.0	171174.0	2370
2010	98698.0	21101.0	24204.0	158850.0	1981

```
[34]: df.groupby('month').sum()
```

```
[34]:
```

	Global_active_power	Global_reactive_power	Voltage	\
month				
1	11008.308	947.904	2006703.97	
2	8152.446	821.946	1828442.49	
3	9745.836	988.560	2018621.08	
4	8474.336	1001.432	1956011.97	
5	8809.694	1078.724	2002497.07	
6	7884.870	1171.386	1959560.14	
7	6921.870	1138.596	2034521.85	
8	6927.570	1128.382	1983726.10	
9	8492.082	1049.824	1947568.65	
10	9680.734	1017.446	2097908.71	
11	10113.624	915.628	1928346.74	
12	11223.876	968.512	1994222.58	

	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3	\
month					
1	46386.6	11270.0	10824.0	61548.0	
2	34349.2	6770.0	9704.0	46932.0	
3	41139.8	10014.0	13379.0	58749.0	
4	35922.0	8530.0	10991.0	50916.0	
5	37616.8	10187.0	11647.0	55326.0	
6	33819.6	9297.0	10440.0	49966.0	
7	29844.6	6263.0	9510.0	41610.0	
8	29753.4	7099.0	9424.0	44356.0	
9	36132.8	10411.0	10226.0	51863.0	
10	40851.6	9198.0	10639.0	55897.0	
11	42647.4	10149.0	10803.0	55449.0	
12	47097.2	10651.0	11203.0	61532.0	

	year
month	
1	245034

```

2      224950
3      247044
4      237003
5      247044
6      234992
7      247044
8      239004
9      234990
10     247044
11     230970
12     240950

```

Drop year and month column. We have created these two just to perform basic analysis

```
[35]: df.shape
```

```
[35]: (1432, 10)
```

```
[36]: df.duplicated().sum()
```

```
[36]: 0
```

Remove year- 2006 because it may create problem while analysis

```
[37]: df = df[df['Date']>'2006-12-31']
```

```
[38]: df.shape
```

```
[38]: (1416, 10)
```

```
[39]: df.isnull().sum()
```

```

[39]: Date                0
      Global_active_power  0
      Global_reactive_power  0
      Voltage              0
      Global_intensity      0
      Sub_metering_1         0
      Sub_metering_2         0
      Sub_metering_3         0
      year                  0
      month                  0
      dtype: int64

```

- Global active power is the real power consumption i.e. the power consumed by electrical appliances other than the sub metered appliances.
- Active energy consumed every minute (in watt hour) = $\frac{\text{Global_active_power} \times 1000}{60} - \text{Sub_metering_1} - \text{Sub_metering_2} - \text{Sub_metering_3}$

- Active energy consumed every minute (in watt hour) represents the active energy consumed every minute (in watt hour) in the household by electrical equipment not measured in sub-meterings 1, 2 and 3

Calculation of target variable - "power_consumption"

```
[40]: a = (df['Global_active_power']*1000/60)
      b = df['Sub_metering_1'] + df['Sub_metering_2'] + df['Sub_metering_3']
      df['power_consumption'] = a - b
      df.head()
```

```
[40]:      Date  Global_active_power  Global_reactive_power  Voltage  \
16 2007-01-01          136.092             8.790  17063.52
17 2007-01-02           96.230             7.922  20438.87
18 2007-01-03           25.416             4.938  17871.77
19 2007-01-04           97.752             8.124  14654.69
20 2007-01-05           72.016             7.874  15283.45

      Global_intensity  Sub_metering_1  Sub_metering_2  Sub_metering_3  year  \
16              563.4             0.0             17.0            209.0  2007
17              403.2             8.0             14.0            681.0  2007
18              105.0             0.0             15.0             89.0  2007
19              410.4            88.0            130.0            795.0  2007
20              316.2           338.0             28.0            300.0  2007

      month  power_consumption
16      1          2042.200000
17      1           900.833333
18      1           319.600000
19      1           616.200000
20      1           534.266667
```

Sum all the values of sub meters into one features i.e., "Sub_metering"

```
[41]: df['Sub_metering']=df['Sub_metering_1']+df['Sub_metering_2']+df['Sub_metering_3']
```

```
[42]: df = df.drop(['Sub_metering_1','Sub_metering_2','Sub_metering_3'],axis=1)
```

```
[43]: df.head()
```

```
[43]:      Date  Global_active_power  Global_reactive_power  Voltage  \
16 2007-01-01          136.092             8.790  17063.52
17 2007-01-02           96.230             7.922  20438.87
18 2007-01-03           25.416             4.938  17871.77
19 2007-01-04           97.752             8.124  14654.69
20 2007-01-05           72.016             7.874  15283.45

      Global_intensity  year  month  power_consumption  Sub_metering
```

16	563.4	2007	1	2042.200000	226.0
17	403.2	2007	1	900.833333	703.0
18	105.0	2007	1	319.600000	104.0
19	410.4	2007	1	616.200000	1013.0
20	316.2	2007	1	534.266667	666.0

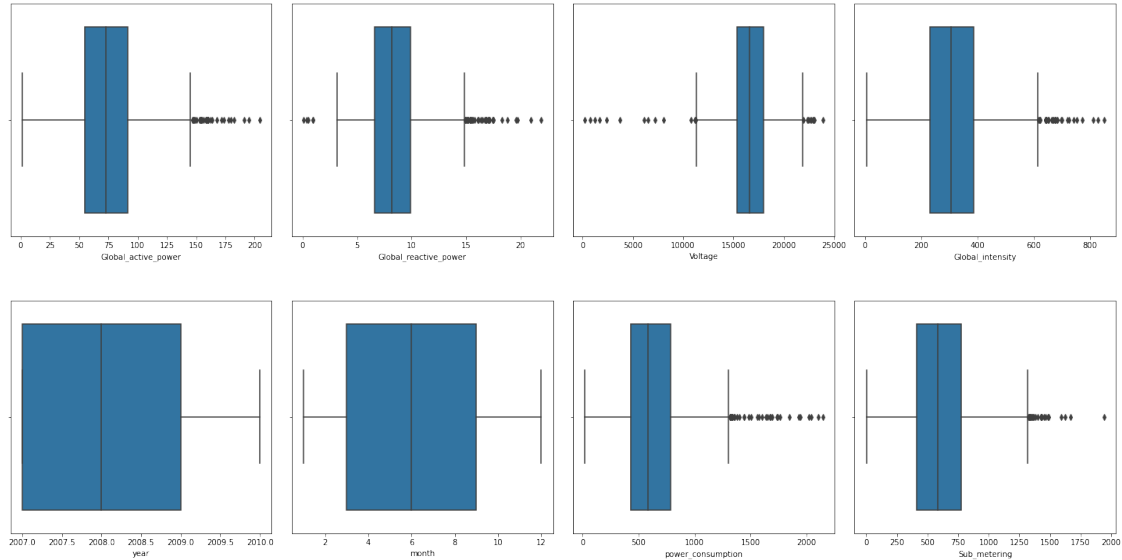
```
[44]: df.isnull().sum()
```

```
[44]: Date                                0
      Global_active_power                 0
      Global_reactive_power              0
      Voltage                             0
      Global_intensity                   0
      year                               0
      month                              0
      power_consumption                  0
      Sub_metering                       0
      dtype: int64
```

3 Graphical Analysis

3.1 Outliers

```
[45]: fig, ax = plt.subplots(ncols=4, nrows=2, figsize=(20,10))
      index = 0
      ax = ax.flatten()
      for col, value in df.items():
          if col!='Date':
              sns.boxplot(x = col, data = df, ax=ax[index])
              index += 1
      plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



```
[46]: def find_boundaries(df, variable, distance):
    IQR = df[variable].quantile(0.75) - df[variable].quantile(0.25)
    lower_boundary = df[variable].quantile(0.25) - (IQR*distance)
    upper_boundary = df[variable].quantile(0.75) + (IQR*distance)
    return upper_boundary, lower_boundary
```

```
[47]: outliers_columns = ['Global_active_power',
    ↪ 'Global_reactive_power', 'Voltage', 'Global_intensity', 'power_consumption', 'Sub_metering']
for i in outliers_columns:
    upper_boundary, lower_boundary = find_boundaries(df, i, 1.5)
    outliers = np.where(df[i] > upper_boundary, True, np.where(df[i] <
    ↪ lower_boundary, True, False))
    outliers_df = df.loc[outliers, i]
    df_trimed = df.loc[~outliers, i]
    df[i] = df_trimed
```

```
[48]: df.isnull().sum()
```

```
[48]: Date                0
Global_active_power      33
Global_reactive_power    42
Voltage                  23
Global_intensity         32
year                    0
month                   0
power_consumption       34
Sub_metering            24
dtype: int64
```

```
[49]: df.fillna(df.median().round(1), inplace=True)
```

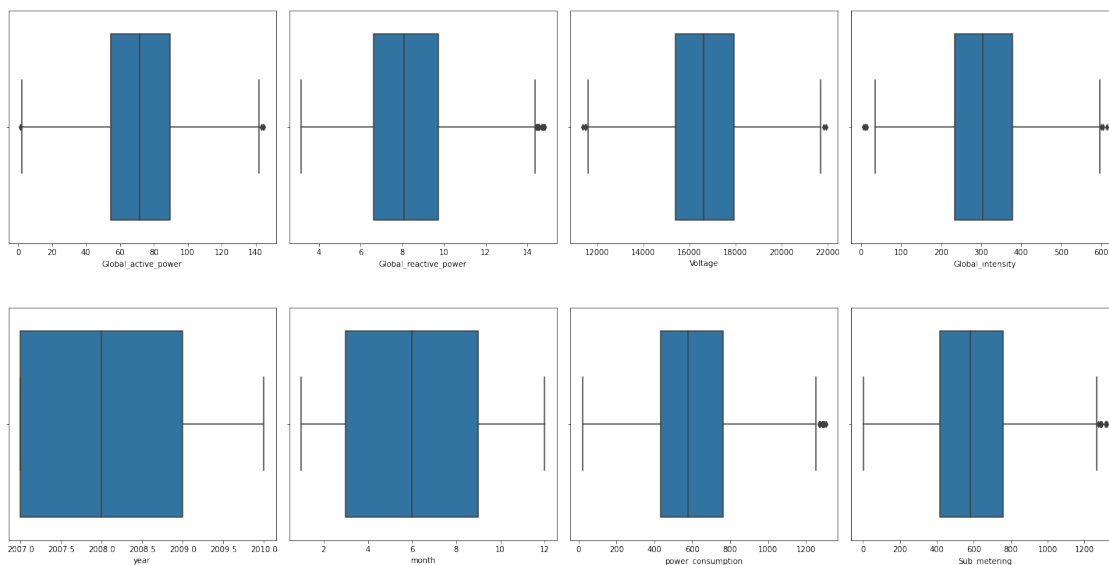
```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning:
DataFrame.mean and DataFrame.median with numeric_only=None will include
datetime64 and datetime64tz columns in a future version.
    """Entry point for launching an IPython kernel.
```

```
[50]: df.dropna(inplace = True)
```

```
[51]: df.isnull().sum()
```


```
[51]: Date                                0
Global_active_power                      0
Global_reactive_power                    0
Voltage                                  0
Global_intensity                         0
year                                     0
month                                    0
power_consumption                       0
Sub_metering                            0
dtype: int64
```

```
[52]: fig, ax = plt.subplots(ncols=4, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    if col!='Date':
        sns.boxplot(x = col, data = df, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```




3.2 Barplot

```
[53]: fig, ax = plt.subplots(ncols = 5, nrows = 1, figsize=(20,5))
index = 0
ax = ax.flatten()
for col, value in df.items():
    if col not in ['Date', 'year', 'month']:
        sns.barplot(y = df[col], x = df['year'], data = df, ax=ax[index] )
        index += 1
    if index == 5:
        break
plt.tight_layout(pad=1, w_pad=1, h_pad=10.0)
```



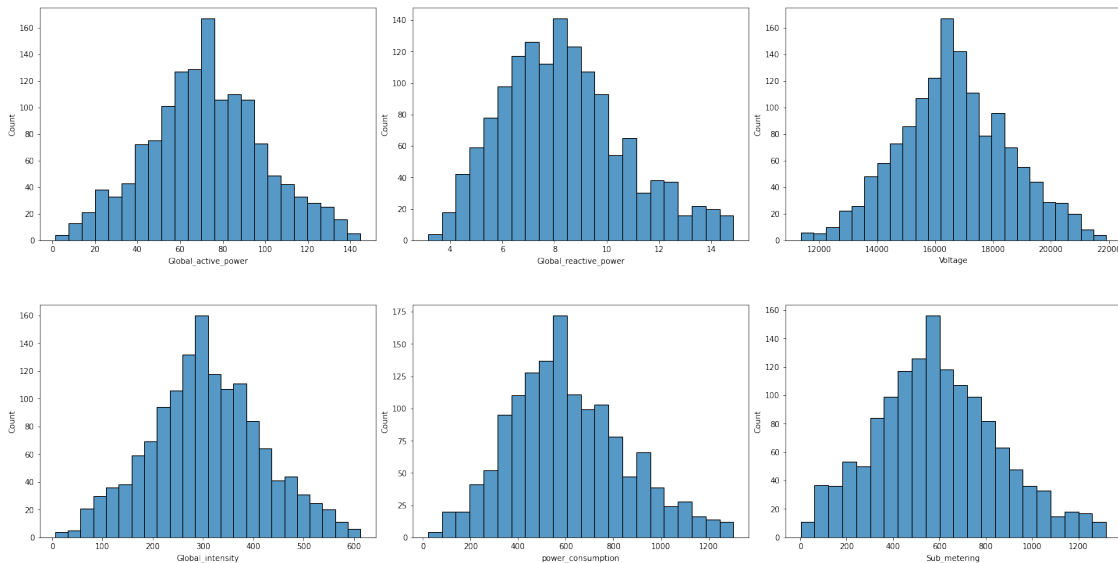
```
[54]: fig, ax = plt.subplots(ncols = 5, nrows = 1, figsize=(20,5))
index = 0
ax = ax.flatten()
for col, value in df.items():
    if col not in ['Date', 'year', 'month']:
        sns.barplot(y = df[col], x = df['month'], data = df, ax=ax[index] )
        index += 1
    if index == 5:
        break
plt.tight_layout(pad=1, w_pad=1, h_pad=10.0)
```



Observation * Power consumption in November and January are on higher side * Voltage is almost equal in every month

3.3 Histplot

```
[55]: fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    if col not in ['Date', 'year', 'month']:
        sns.histplot(value, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



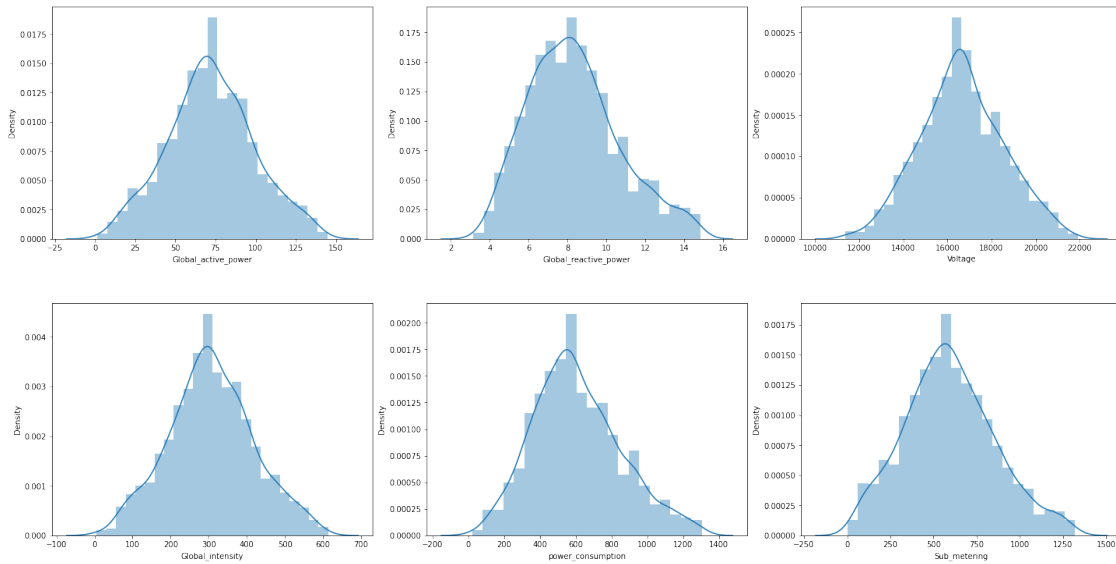
3.4 Distplot

```
[56]: fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    if col not in ['Date', 'year', 'month']:
        sns.distplot(value, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```

```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
    warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
    warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
    warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
    warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
    warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
    warnings.warn(msg, FutureWarning)

```



4 Statistical Analysis

```
[57]: df.corr()
```

```
[57]:
```

	Global_active_power	Global_reactive_power	Voltage	\
Global_active_power	1.000000	0.078498	0.248528	
Global_reactive_power	0.078498	1.000000	0.329612	
Voltage	0.248528	0.329612	1.000000	
Global_intensity	0.996324	0.099367	0.253729	
year	0.015729	0.113831	0.020252	
month	0.024475	0.044594	0.056966	
power_consumption	0.819335	-0.039454	0.227614	
Sub_metering	0.827409	0.196993	0.208694	


	Global_intensity	year	month	\
Global_active_power	0.996324	0.015729	0.024475	
Global_reactive_power	0.099367	0.113831	0.044594	
Voltage	0.253729	0.020252	0.056966	
Global_intensity	1.000000	0.005334	0.021361	
year	0.005334	1.000000	-0.036917	
month	0.021361	-0.036917	1.000000	
power_consumption	0.818675	-0.067469	0.039837	
Sub_metering	0.832383	0.069651	0.002115	

	power_consumption	Sub_metering
Global_active_power	0.819335	0.827409

Global_reactive_power	-0.039454	0.196993
Voltage	0.227614	0.208694
Global_intensity	0.818675	0.832383
year	-0.067469	0.069651
month	0.039837	0.002115
power_consumption	1.000000	0.520581
Sub_metering	0.520581	1.000000

```
[58]: sns.set(rc={'figure.figsize':(15,10)})
sns.heatmap(data=df.corr(), annot=True, vmin=-1, vmax=1)
```

```
[58]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc9c7e72810>
```



Observation * Global_active_power, Global_intensity and sub_metering are highly correlated

```
[59]: df.describe().T
```

	count	mean	std	min	\
Global_active_power	1416.0	72.277684	27.051160	1.372000	
Global_reactive_power	1416.0	8.335674	2.362348	3.160000	
Voltage	1416.0	16657.260410	1895.779688	11361.000000	

Global_intensity	1416.0	306.822458	112.187710	5.800000
year	1416.0	2008.455508	1.104183	2007.000000
month	1416.0	6.435734	3.416252	1.000000
power_consumption	1416.0	607.389736	245.578489	20.866667
Sub_metering	1416.0	593.437853	262.875053	2.000000

	25%	50%	75%	max
Global_active_power	54.5390	71.6	89.7280	144.784000
Global_reactive_power	6.6215	8.1	9.7285	14.832000
Voltage	15396.3875	16618.0	17948.9475	21911.440000
Global_intensity	233.1500	304.0	379.1000	613.600000
year	2007.0000	2008.0	2009.0000	2010.000000
month	3.0000	6.0	9.0000	12.000000
power_consumption	432.9750	579.6	761.0750	1304.166667
Sub_metering	415.0000	579.5	757.2500	1320.000000

Observation * Maximum power consumption in a day is 2146 w/h * Average consumption is 631 w/h * Minimum consumption is 21 w/h

5 Segregating Independent and Dependent Features

```
[60]: X = df.iloc[ : , [1,2,3,4,6,8]]
      y = df.iloc[ : , -2]
```

```
[61]: X.shape
```

```
[61]: (1416, 6)
```

```
[62]: y.shape
```

```
[62]: (1416,)
```

```
[63]: X.head()
```

```
[63]:   Global_active_power  Global_reactive_power  Voltage  Global_intensity \
16              136.092              8.790  17063.52              563.4
17              96.230              7.922  20438.87              403.2
18              25.416              4.938  17871.77              105.0
19              97.752              8.124  14654.69              410.4
20              72.016              7.874  15283.45              316.2

      month  Sub_metering
16        1          226.0
17        1          703.0
18        1          104.0
```

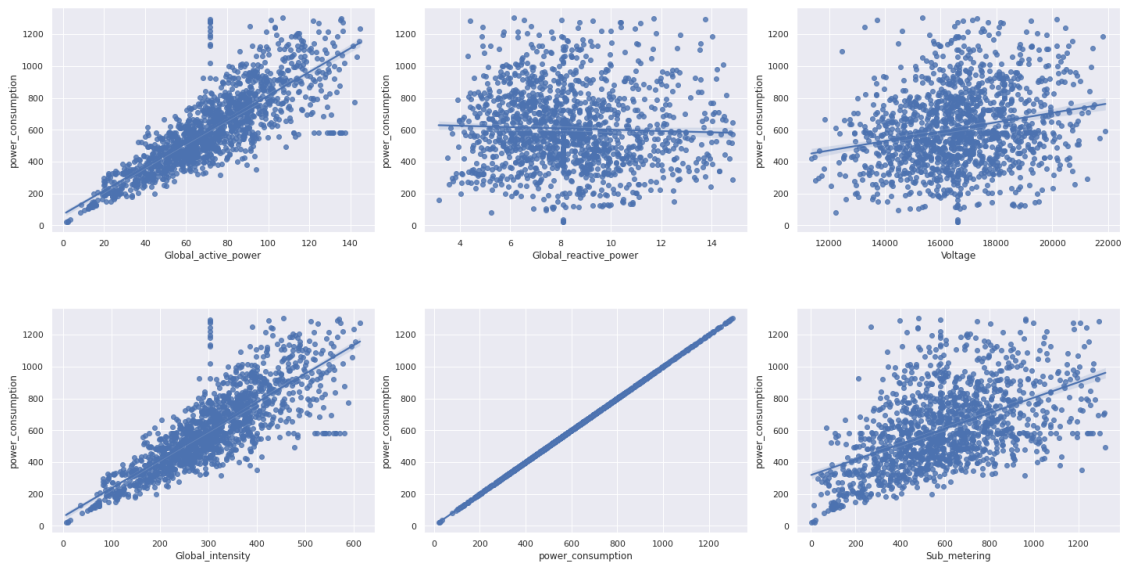
19	1	1013.0
20	1	666.0

```
[64]: y.head()
```

```
[64]: 16    579.600000
      17    900.833333
      18    319.600000
      19    616.200000
      20    534.266667
      Name: power_consumption, dtype: float64
```

Regplot

```
[65]: fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
      index = 0
      ax = ax.flatten()
      for col, value in df.items():
          if col not in ['Date', 'year', 'month']:
              sns.regplot(x = df[col], y = df["power_consumption"], data = df , ax = ax[
                  ↳ax[index])
                  index += 1
      plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



6 Train Test Split

```
[66]: from sklearn.model_selection import train_test_split

[67]: X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=7,test_size=0.
      ↪ 33)
```

7 Scaling

- Not required for Decision Trees

```
[68]: # from sklearn.preprocessing import StandardScaler

[69]: # scaler=StandardScaler()

[70]: # X_train = scaler.fit_transform(X_train)

[71]: # X_test = scaler.transform(X_test)

[72]: # len(X_train)

[73]: # mm
```

8 Save Preprocess Model Data Using Pickle

```
[74]: # preprocess_model = [X_train,y_train,X_test,y_test]

[75]: # import pickle

[76]: # pickle.dump(preprocess_model, open('preprocess_model.pkl','wb'))

[77]: # preprocess_model = pickle.load(open('preprocess_model.pkl','rb'))
```

Note * We have successfully stored our scaled data into pickel file so we can use it further in other file by just importing it

9 Save Data into MongoDB

```
[78]: # y_train.T
```

```
[79]: # database_df = pd.DataFrame([X_train.T[0],X_train.T[1],X_train.T[2],X_train.
      ↪T[3], X_train.T[4], X_train.T[5],y_train]).T

[80]: # database_df.columns=['Global_active_power', 'Global_reactive_power',␣
      ↪'Voltage', 'Global_intensity', 'month', 'Sub_metering', 'power_consumption']

[81]: # database_df.head()

[82]: # l=[]
      # for i ,row in database_df.iterrows():
      #     l.append(dict(row))

[83]: # import pymongo
      # from pymongo import MongoClient

[84]: # client = pymongo.MongoClient("mongodb+srv://subhashdixit17:Anushka27@cluster0.
      ↪elq8eyt.mongodb.net/?retryWrites=true&w=majority")

[85]: # db=client['Household_Power_Preprocessed_Data']
      # collections = db['Training__Independent_and_Dependent_Dataset']
      # collections.insert_many(l)
```

10 Load Preprocessed data using Pickle

```
[86]: # preprocess_model = pickle.load(open('preprocess_model.pkl','rb'))

[87]: # X_train = preprocess_model[0]
      # y_train = preprocess_model[1]
      # X_test = preprocess_model[2]
      # y_test = preprocess_model[3]

[88]: # X_train =pd.DataFrame(X_train)
      # X_test =pd.DataFrame(X_test)
      # X_train.columns=['Global_active_power', 'Global_reactive_power', 'Voltage',␣
      ↪'Global_intensity', 'month', 'Sub_metering']
      # X_test.columns=['Global_active_power', 'Global_reactive_power', 'Voltage',␣
      ↪'Global_intensity', 'month', 'Sub_metering']
```

11 VIF Check

- To check multicollinearity

```
[89]: # X_train2 = X_train.copy()
# X_train= pd.DataFrame(X_train)
```

```
[90]: # X_train
```

```
[91]: # from statsmodels.stats.outliers_influence import variance_inflation_factor
# vif = [variance_inflation_factor(X_train.values, i) for i in range(X_train.
#       ↪shape[1])]
# print(X_train.columns)
# print(vif)
```

```
[92]: # while (max(vif) > 5):
#     indx = vif.index(max(vif)) #Get the index of variable with highest VIF
#     print(indx)
#     X_train.drop(X_train.columns[indx],axis = 1, inplace = True)
#     vif = [variance_inflation_factor(X_train.values, i) for i in
#           ↪range(X_train.shape[1])]
# vif = [variance_inflation_factor(X_train.values, i) for i in range(X_train.
#       ↪shape[1])]
# print(X_train.columns)
# print(vif)
```

```
[93]: # X_test = pd.DataFrame(X_test)
# X_test = X_test[X_train.columns]
```

12 Model Creation

```
[94]: from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
from sklearn.tree import DecisionTreeRegressor
from sklearn.tree import ExtraTreeRegressor
from sklearn.model_selection import GridSearchCV
```

```
[95]: parameters = {
    "splitter":["best","random"],
    "max_depth" : [1,3,5,7,9,11,12],
    "min_samples_leaf":[1,2,3,4,5,6,7,8,9,10],
    "min_weight_fraction_leaf":[0.1,0.2,0.3,0.4,0.5],
    "max_features":["auto","log2","sqrt",None],
    "max_leaf_nodes":[None,10,20,30,40,50,60,70,80,90]
}

## We will train that models
models = {
    1: DecisionTreeRegressor(random_state=0),
```

```

2: ExtraTreeRegressor(random_state=0),
3: GridSearchCV(DecisionTreeRegressor(random_state=42), param_grid=parameters,
↳ verbose=1, cv=3),
4: GridSearchCV(ExtraTreeRegressor(random_state=42), param_grid=parameters,
↳ verbose=1, cv=3)
}

```

```
[96]: map_keys = list(models.keys())
```

```
[97]: # Get model name using id from linear_model_collection
def get_model_building_technique_name(num):
    if num == 1:
        return 'DecisionTreeRegressor()'
    if num == 2:
        return 'ExtraTreeRegressor()'
    if num == 3:
        return "GridSearchCV()_DTR"
    if num == 4:
        return "GridSearchCV()_ETR"
    return ''

```

```
[98]: results = [];
for key_index in range(len(map_keys)):
    key = map_keys[key_index]
    model = models[key]
    print(key_index)
    model.fit(X_train, y_train)

    '''Test Accuracy'''
    y_pred = model.predict(pd.DataFrame(X_test))

    R_Squared_Test = r2_score(y_test, y_pred)
    Adjusted_R_Squared_Test = (1 - (1-R_Squared_Test)*(len(y_test)-1)/
↳ (len(y_test)-X_test.shape[1]-1))

    '''Train Accuracy'''
    y_pred_train = model.predict(X_train)

    R_Squared_Train = r2_score(y_train, y_pred_train)
    Adjusted_R_Squared_Train = (1 - (1-R_Squared_Train)*(len(y_train)-1)/
↳ (len(y_train)-X_test.shape[1]-1))

    results.append({
        'Model Name' : get_model_building_technique_name(key),
        'Trained Model' : model,
        'R_Squared_Test' : R_Squared_Test,

```

```

    'Adjusted_R_Squared_Test' : Adjusted_R_Squared_Test,
    'R_Squared_Train' : R_Squared_Train,
    'Adjusted_R_Squared_Train' : Adjusted_R_Squared_Train
})

```

0

1

2

Fitting 3 folds for each of 28000 candidates, totalling 84000 fits

3

Fitting 3 folds for each of 28000 candidates, totalling 84000 fits

12.1 Train and Test Accuracy

```

[99]: result_df = pd.DataFrame(results)
      result_df

```

```

[99]:
      Model Name                                Trained Model \
0  DecisionTreeRegressor()  DecisionTreeRegressor(random_state=0)
1    ExtraTreeRegressor()    ExtraTreeRegressor(random_state=0)
2    GridSearchCV()_DTR  GridSearchCV(cv=3, estimator=DecisionTreeRegre...
3    GridSearchCV()_ETR  GridSearchCV(cv=3, estimator=ExtraTreeRegresso...

```

	R_Squared_Test	Adjusted_R_Squared_Test	R_Squared_Train	\
0	0.879728	0.878163	1.00000	
1	0.858121	0.856275	1.00000	
2	0.712480	0.708737	0.72561	
3	0.712480	0.708737	0.72561	

	Adjusted_R_Squared_Train
0	1.00000
1	1.00000
2	0.72386
3	0.72386

12.2 Checking Best Model

```

[100]: Best_Model_Name = 'GridSearchCV()_DTR' # Because Train and Test accuracy are in
      ↳ range. It means our model is not overfitted.

```


12.3 Save Best Model

```
[102]: import pickle
Best_Trained_model = Best_Model_Name
with open('DTR_Household_Power_Consumption_Regression_Problem.sav', 'wb') as f:
    pickle.dump(Best_Trained_model, f)
```

**

The End

**