

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

In [2]:

```
df=pd.read_csv("adult.csv",na_values='?')
```

In [3]:

```
df.head()
```

Out[3]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	NaN	77053	HS-grad	9	Widowed	NaN	Not-in-family
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family
2	66	NaN	186061	Some-college	10	Widowed	NaN	Unmarried
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child

In [4]:

```
df.shape
```

Out[4]:

```
(32561, 15)
```

In [5]:

```
df.isnull().sum()
```

Out[5]:

```
age                0
workclass          1836
fnlwgt             0
education          0
education.num      0
marital.status     0
occupation         1843
relationship       0
race               0
sex               0
capital.gain       0
capital.loss       0
hours.per.week     0
native.country     583
income             0
dtype: int64
```

In [6]:

```
df.isin(['?']).sum()
```

Out[6]:

```
age                0
workclass          0
fnlwgt             0
education          0
education.num      0
marital.status     0
occupation         0
relationship       0
race               0
sex               0
capital.gain       0
capital.loss       0
hours.per.week     0
native.country     0
income             0
dtype: int64
```

In [7]:

```
df['workclass'].value_counts()
```

Out[7]:

Private	22696
Self-emp-not-inc	2541
Local-gov	2093
State-gov	1298
Self-emp-inc	1116
Federal-gov	960
Without-pay	14
Never-worked	7

Name: workclass, dtype: int64

In [8]:

```
df['occupation'].value_counts()
```

Out[8]:

Prof-specialty	4140
Craft-repair	4099
Exec-managerial	4066
Adm-clerical	3770
Sales	3650
Other-service	3295
Machine-op-inspct	2002
Transport-moving	1597
Handlers-cleaners	1370
Farming-fishing	994
Tech-support	928
Protective-serv	649
Priv-house-serv	149
Armed-Forces	9

Name: occupation, dtype: int64

In [9]:

```
plt.figure(figsize=(10,6))
sns.heatmap(df.isna().transpose(),
            cmap="crest")
```

Out[9]:

&lt;AxesSubplot:&gt;



In [10]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   32561 non-null  int64
1   workclass             30725 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education.num        32561 non-null  int64
5   marital.status       32561 non-null  object
6   occupation            30718 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week        32561 non-null  int64
13  native.country       31978 non-null  object
14  income               32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

In [11]:

df.columns

Out[11]:

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education.num',
      'marital.status', 'occupation', 'relationship', 'race', 'sex',
      'capital.gain', 'capital.loss', 'hours.per.week', 'native.country',
      'income'],
      dtype='object')
```

In [12]:

df.describe().T

Out[12]:

	count	mean	std	min	25%	50%	75%
<b>age</b>	32561.0	38.581647	13.640433	17.0	28.0	37.0	48.0
<b>fnlwgt</b>	32561.0	189778.366512	105549.977697	12285.0	117827.0	178356.0	237051.0
<b>education.num</b>	32561.0	10.080679	2.572720	1.0	9.0	10.0	12.0
<b>capital.gain</b>	32561.0	1077.648844	7385.292085	0.0	0.0	0.0	0.0
<b>capital.loss</b>	32561.0	87.303830	402.960219	0.0	0.0	0.0	0.0
<b>hours.per.week</b>	32561.0	40.437456	12.347429	1.0	40.0	40.0	45.0

In [13]:

```
numerical_feature=[i for i in df.columns if df[i].dtype!='O']  
categorical_feature=[i for i in df.columns if df[i].dtype=='O']
```

In [14]:

```
numerical_feature
```

Out[14]:

```
['age',  
 'fnlwgt',  
 'education.num',  
 'capital.gain',  
 'capital.loss',  
 'hours.per.week']
```

In [15]:

```
categorical_feature
```

Out[15]:

```
['workclass',  
 'education',  
 'marital.status',  
 'occupation',  
 'relationship',  
 'race',  
 'sex',  
 'native.country',  
 'income']
```

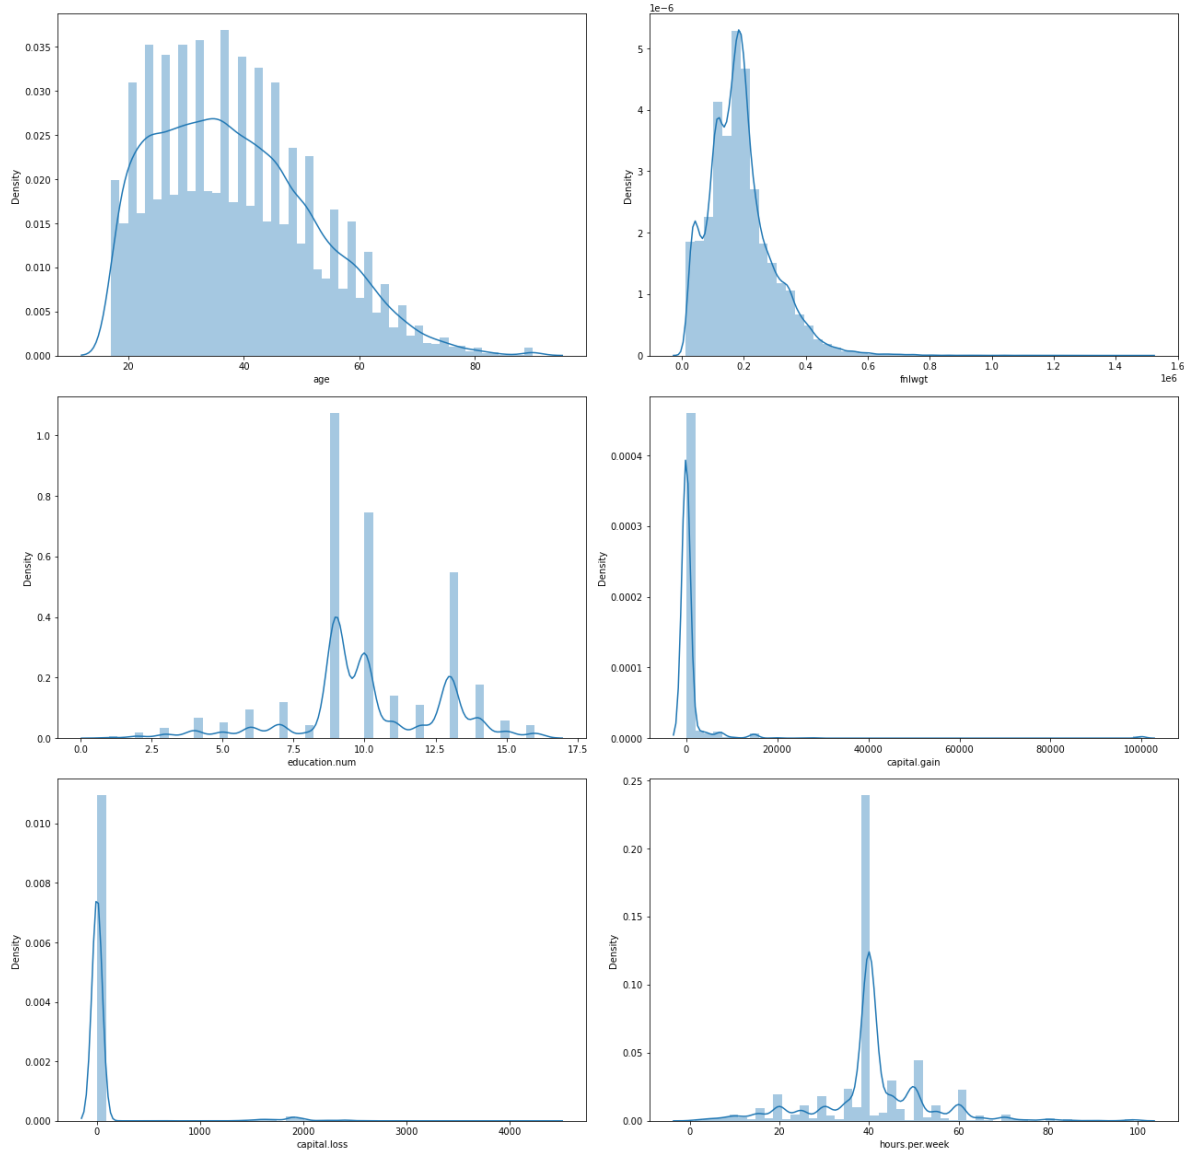
## Univariate analysis

In [16]:

```
plt.figure(figsize=(18,18))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20, fontweight='bold', a

for i in range(0, len(numerical_feature)):
    plt.subplot(3, 2, i+1)
    sns.distplot(df[numerical_feature[i]],label='income')
    plt.xlabel(numerical_feature[i])
    plt.tight_layout()
```

Univariate Analysis of Numerical Features

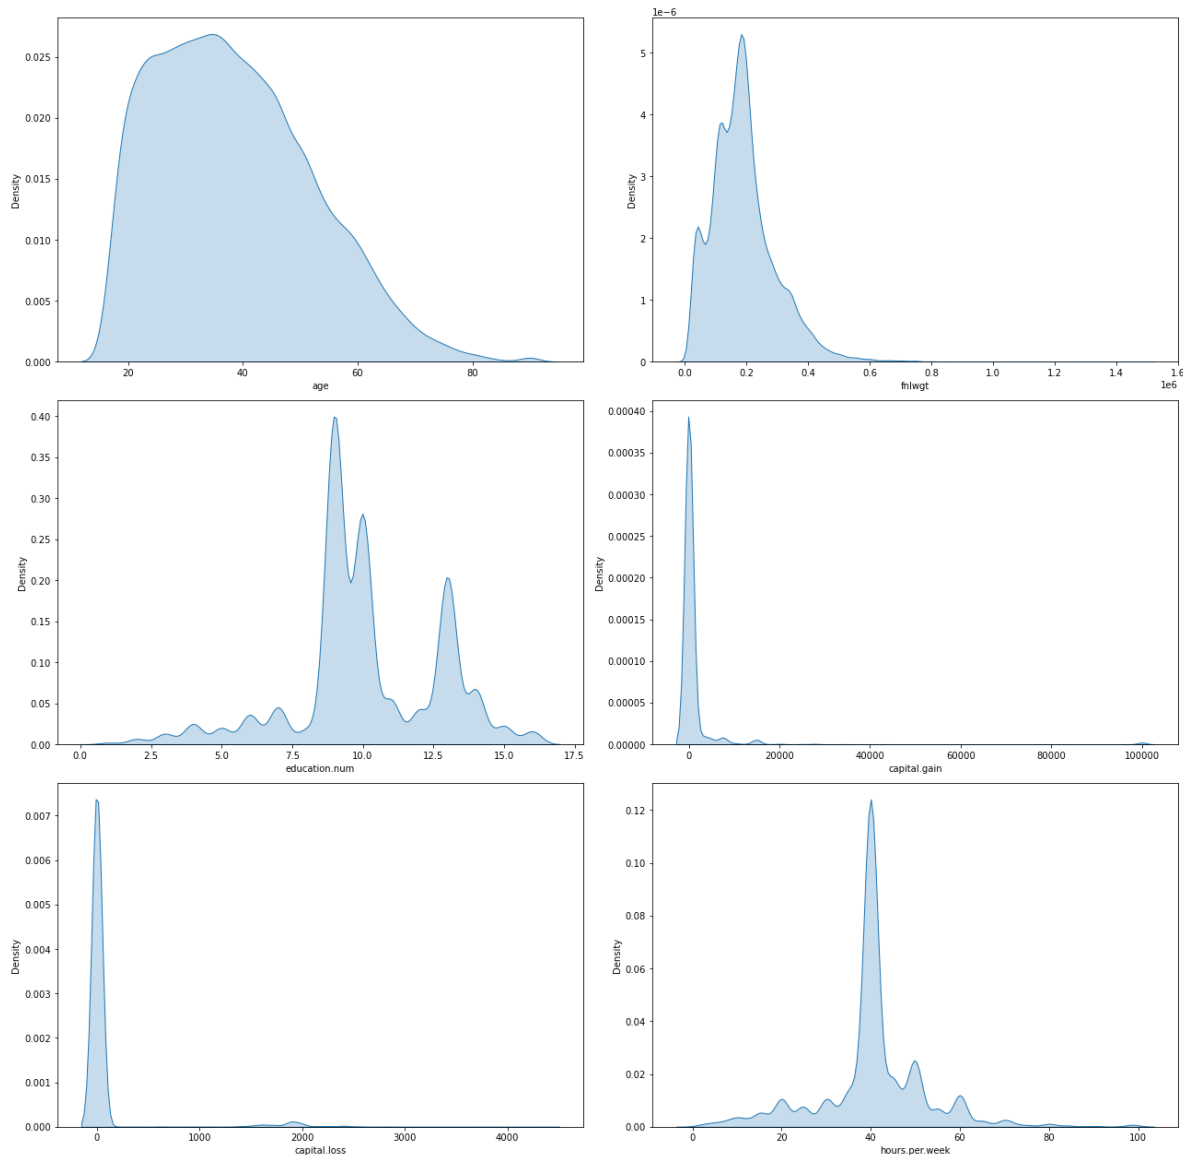


In [17]:

```
plt.figure(figsize=(18,18))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20, fontweight='bold', a

for i in range(0, len(numerical_feature)):
    plt.subplot(3, 2, i+1)
    sns.kdeplot(x=df[numerical_feature[i]],shade=True,data=df)
    plt.xlabel(numerical_feature[i])
    plt.tight_layout()
```

Univariate Analysis of Numerical Features





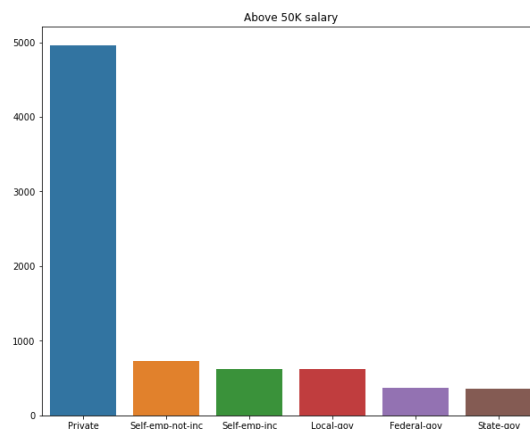

In [18]:

```
def cate_feature(df,categorical_feature):
    plt.figure(figsize=(22,8))
    a=df[df['income']=='<=50K'][categorical_feature].value_counts() # Below 50k salary
    b=df[df['income']=='>50K'][categorical_feature].value_counts() # Above 50k salary
    plt.subplot(1,2,1)
    plt.title('Below 50K salary')
    sns.barplot(a.index,a.values)
    plt.subplot(1,2,2)
    plt.title('Above 50K salary')
    sns.barplot(b.index,b.values)
    plt.show()
```

In [19]:

```
for i in categorical_feature:
    print(i,':\n')
    cate_feature(df,i)
```

workclass :



education :

Below 50K salary

Above 50K salary

```
plt.figure(figsize=(18,18))
plt.suptitle('Univariate Analysis of Categorical Features', fontsize=20, fontweight='bold',
for i in range(0, len(categorical_feature)):
    plt.subplot(5, 2, i+1)
    sns.countplot(x=df[categorical_feature[i]],data=df)
    plt.xlabel(categorical_feature[i])
    plt.tight_layout()
```




In [21]:

```
plt.figure(figsize=(16,6))  
sns.countplot(x="education",hue="income",data=df)
```

Out[21]:

<AxesSubplot:xlabel='education', ylabel='count'>

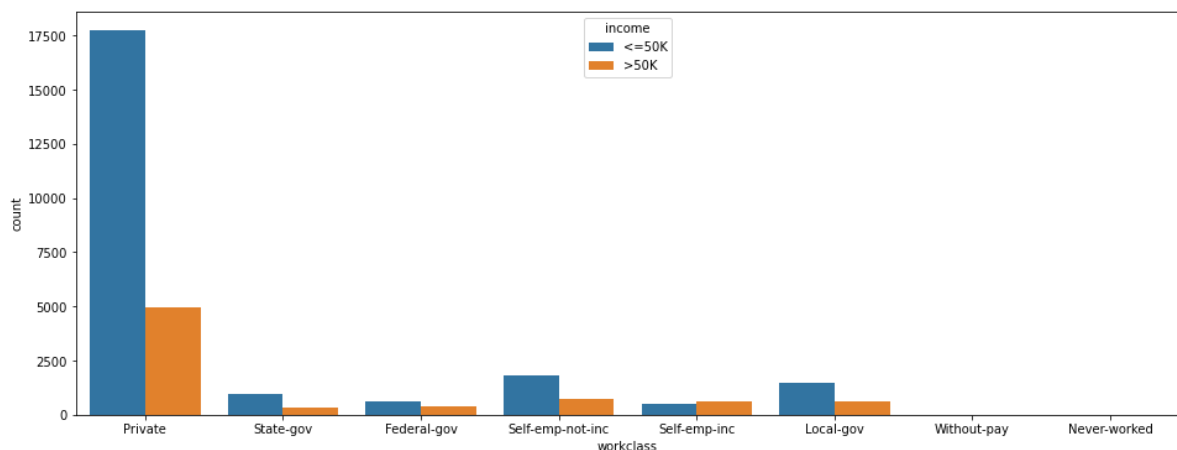


In [22]:

```
plt.figure(figsize=(16,6))  
sns.countplot(x="workclass",hue="income",data=df)
```

Out[22]:

<AxesSubplot:xlabel='workclass', ylabel='count'>




In [23]:

```
plt.figure(figsize=(15,8))
sns.set_theme(style="darkgrid")
sns.lineplot(x='age', y='hours.per.week',
             hue="income",
             data=df)
```


Out[23]:

&lt;AxesSubplot:xlabel='age', ylabel='hours.per.week'&gt;



In [24]:

```
plt.figure(figsize=(15, 10))
sns.heatmap(df.corr(), annot=True, cmap="PuOr", annot_kws={"size": 10})
plt.show()
```



## Feature Engineering

In [25]:


```
df=df.dropna(axis=0)
```

In [26]:

```
plt.figure(figsize=(10,6))
sns.heatmap(df.isna().transpose(),
            cmap="PiYG")
```

Out[26]:

&lt;AxesSubplot:&gt;



## Label Encoding for Categorical feature

In [27]:

```
from sklearn.preprocessing import LabelEncoder
```

In [28]:

```
le=LabelEncoder()
```

In [29]:

```
for i in categorical_feature :
    df[i] = le.fit_transform(df[i])
```

In [30]:

```
df=df.drop(['marital.status','relationship','capital.gain','capital.loss'],axis=1)
```

In [31]:

```
df.head()
```

Out[31]:

	age	workclass	fnlwgt	education	education.num	occupation	race	sex	hours.per.week
1	82	2	132870	11	9	3	4	0	18
3	54	2	140359	5	4	6	4	0	40
4	41	2	264663	15	10	9	4	0	40
5	34	2	216864	11	9	7	4	0	45
6	38	2	150601	0	6	0	4	1	40

In [32]:

```
X=df.drop(['income'],axis=1)  
y=df['income']
```

In [33]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=100)
```

In [34]:

```
X_train.shape
```

Out[34]:

```
(20208, 10)
```

In [35]:

```
y_train.shape
```

Out[35]:

```
(20208,)
```

In [36]:

```
X_test.shape
```

Out[36]:

```
(9954, 10)
```

In [37]:

```
y_test.shape
```

Out[37]:

```
(9954,)
```

In [38]:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import ExtraTreesClassifier
```

## Decision Tree Classifier

In [39]:

```
dc=DecisionTreeClassifier()
```

In [40]:

```
dc.fit(X_train,y_train)
```

Out[40]:

```
▼ DecisionTreeClassifier
DecisionTreeClassifier()
```

In [41]:

```
dc.score(X_train,y_train)
```

Out[41]:

```
0.9998020585906572
```

In [42]:

```
dc_pred=dc.predict(X_test)
```

In [43]:

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.model_selection import cross_val_score
```

In [44]:

```
accuracy_score(y_test,dc_pred)
```

Out[44]:

```
0.7458308217801889
```



In [45]:

```
print(classification_report(y_test,dc_pred))
```

	precision	recall	f1-score	support
0	0.83	0.83	0.83	7458
1	0.49	0.50	0.50	2496
accuracy			0.75	9954
macro avg	0.66	0.66	0.66	9954
weighted avg	0.75	0.75	0.75	9954

In [46]:

```
grid_param_dc = {
    'criterion': ['gini', 'entropy'],
    'max_depth' : range(2,10,1),
    'min_samples_leaf' : range(1,8,1),
    'min_samples_split': range(2,8,1),
    'splitter' : ['best', 'random']
}
```

In [47]:


```
from sklearn.model_selection import GridSearchCV
grid_search_dc=GridSearchCV(estimator=dc,param_grid=grid_param_dc,cv=3,verbose=1)
```

In [48]:

```
grid_search_dc.fit(X_train,y_train)
```

Fitting 3 folds for each of 1344 candidates, totalling 4032 fits

Out[48]:



In [49]:

```
grid_search_dc.best_params_
```

Out[49]:

```
{'criterion': 'entropy',
 'max_depth': 8,
 'min_samples_leaf': 7,
 'min_samples_split': 3,
 'splitter': 'random'}
```

In [50]:

```
dc_model_with_best_parm=DecisionTreeClassifier(criterion='gini',
max_depth= 5,
min_samples_leaf=6,
min_samples_split= 2,
splitter= 'best')
```

In [51]:

```
dc_model_with_best_parm.fit(X_train,y_train)
```

Out[51]:

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5, min_samples_leaf=6)
```

In [52]:

```
pred_dc_model_with_best_parm=dc_model_with_best_parm.predict(X_test)
```

In [53]:

```
print(accuracy_score(y_test,pred_dc_model_with_best_parm))
```

0.7976692786819369

In [54]:

```
print(classification_report(y_test,pred_dc_model_with_best_parm))
```

	precision	recall	f1-score	support
0	0.82	0.94	0.87	7458
1	0.67	0.38	0.49	2496
accuracy			0.80	9954
macro avg	0.74	0.66	0.68	9954
weighted avg	0.78	0.80	0.78	9954

In [55]:

```
print(confusion_matrix(y_test,pred_dc_model_with_best_parm))
```

```
[[6980  478]
 [1536  960]]
```

## RandomForestClassifier

In [56]:

```
rf= RandomForestClassifier()
```

In [57]:

```
rf.fit(X_train,y_train)
```

Out[57]:

```
▼ RandomForestClassifier  
RandomForestClassifier()
```

In [58]:

```
rf.score(X_train,y_train)
```

Out[58]:

```
0.9997525732383215
```

In [59]:

```
rf_pred=rf.predict(X_test)
```

In [60]:

```
print(accuracy_score(y_test,rf_pred))
```

```
0.7958609604179224
```

In [61]:

```
rf_grid_param = {  
    "n_estimators" : [90,100,110,120],  
    'criterion': ['gini', 'entropy'],  
    'max_depth' : range(2,20,1),  
    'min_samples_leaf' : range(1,10,1),  
    'min_samples_split': range(2,10,1),  
    'max_features' : ['auto','log2']  
}
```

In [62]:

```
from tpot import TPOTClassifier
```

In [63]:

```
rf_tpot_classifier = TPOTClassifier(generations= 5, population_size= 28, offspring_size= 14,  
    verbosity= 2, early_stop= 12,  
    config_dict={'sklearn.ensemble.RandomForestClassifier': rf  
    cv = 5, scoring = 'accuracy'})
```

In [64]:

```
rf_tpot_classifier.fit(X_train,y_train)
```

Generation 1 - Current best internal CV score: 0.8133910201238885

Generation 2 - Current best internal CV score: 0.8133910201238885

Generation 3 - Current best internal CV score: 0.8134405985359614

Generation 4 - Current best internal CV score: 0.8138856042731719

Generation 5 - Current best internal CV score: 0.8138856042731719

Best pipeline: RandomForestClassifier(RandomForestClassifier(input\_matrix, criterion=gini, max\_depth=7, max\_features=auto, min\_samples\_leaf=6, min\_samples\_split=4, n\_estimators=90), criterion=gini, max\_depth=6, max\_features=auto, min\_samples\_leaf=6, min\_samples\_split=5, n\_estimators=110)

Out[64]:

```

TPOTClassifier
TPOTClassifier(config_dict={'sklearn.ensemble.RandomForestClassifier':
{'criterion': ['gini',
'entropy'],
'max_depth': range(2, 20),
'max_features': ['auto',
'log2'],

```

In [65]:

```
rf_accuracy = rf_tpot_classifier.score(X_test, y_test)
print(rf_accuracy)
```

0.8089210367691381

In [66]:

```
rf_g_pred=rf_tpot_classifier.predict(X_test)
```

In [67]:

```
print(accuracy_score(y_test,rf_g_pred))
```

0.8089210367691381

In [68]:

```
#grid_search_rf=GridSearchCV(estimator=rf,param_grid=rf_grid_param,cv=3,verbose=1)
```

In [69]:

```
#grid_search_rf.fit(X_train,y_train)
```

## SVC

In [70]:

```
sv=SVC()
```

In [71]:

```
sv.fit(X_train,y_train)
```

Out[71]:

```
▼ SVC  
SVC()
```

In [72]:

```
sv.score(X_train,y_train)
```

Out[72]:

```
0.7519794140934284
```

In [73]:

```
sv_pred=sv.predict(X_test)
```

In [74]:

```
print(accuracy_score(y_test,sv_pred))
```

```
0.7492465340566606
```

In [75]:

```
sv_grid_param = {  
    "C" : [0.1,1,2,5,10,100],  
    'kernel': ['linear', 'poly', 'rbf', 'sigmoid', 'precomputed'],  
    'gamma':['scale', 'auto']  
}
```

In [76]:

```
#grid_search_svm=GridSearchCV(estimator=sv,param_grid=sv_grid_param,cv=3,verbose=1)
```

In [77]:

```
#grid_search_svm.fit(X_train,y_train)
```

In [78]:

```
#from sklearn.model_selection import RandomizedSearchCV
```

In [79]:

```
#sv_randomcv=RandomizedSearchCV(estimator=sv,param_distributions=sv_grid_param ,n_iter=100,  
                                #random_state=100,n_jobs=-1)
```

In [80]:

```
#sv_randomcv.fit(X_train,y_train)
```

In [81]:

```
#sv_tpot_classifier = TPOTClassifier(generations= 5, population_size= 28, offspring_size= 1  
                                    #verbosity= 2, early_stop= 12,  
                                    #config_dict={'sklearn.svm.SVC': sv_grid_param},  
                                    #cv = 5, scoring = 'accuracy')
```

In [82]:

```
#sv_tpot_classifier.fit(X_train,y_train)
```

## BaggingClassifier

In [83]:

```
bg=BaggingClassifier(n_estimators=10, random_state=0)
```

In [84]:

```
bg.fit(X_train,y_train)
```

Out[84]:

```
▼      BaggingClassifier  
BaggingClassifier(random_state=0)
```

In [85]:

```
bg.score(X_train,y_train)
```

Out[85]:

```
0.9825811559778306
```

In [86]:

```
bg_pred=bg.predict(X_test)
```

In [87]:

```
print(accuracy_score(y_test,bg_pred))
```

0.7858147478400643

In [88]:

```
bg_grid_param = {
    "n_estimators" : [10,20,30,50,100],
    'max_samples': range(1,10,1),
    'oob_score':[True,False]
}
```

In [89]:


```
grid_search_bg=GridSearchCV(estimator=bg,param_grid=bg_grid_param,cv=3,verbose=1)
```

In [90]:

```
grid_search_bg.fit(X_train,y_train)
```

Fitting 3 folds for each of 90 candidates, totalling 270 fits

Out[90]:



In [91]:

```
grid_search_bg.best_params_
```

Out[91]:

```
{'max_samples': 6, 'n_estimators': 10, 'oob_score': True}
```


In [92]:

```
bg_with_best_parm=BaggingClassifier(n_estimators=10,max_samples=6 ,oob_score=True)
```

In [93]:

```
bg_with_best_parm.fit(X_train,y_train)
```

Out[93]:



In [94]:

```
bg_with_best_parm_pred=bg_with_best_parm.predict(X_test)
```

In [95]:

```
print(accuracy_score(y_test,bg_with_best_parm_pred))
```

0.7427164958810528

## ExtraTreesClassifier

In [96]:

```
et=ExtraTreesClassifier()
```

In [97]:

```
et.fit(X_train,y_train)
```

Out[97]:

```
▼ ExtraTreesClassifier  
ExtraTreesClassifier()
```

In [98]:

```
et_pred=et.predict(X_test)
```

In [99]:

```
print(accuracy_score(y_test,et_pred))
```

0.7878239903556359

In [100]:

```
grid_param_et = {  
    'n_estimators':range(100,200,20),  
    'criterion': ['gini','entropy','log_loss'],  
    'max_depth' : range(2,10,1),  
    'min_samples_leaf' : range(1,8,1),  
    'min_samples_split': range(2,8,1),  
    'max_features' : ['sqrt','log2']  
}
```

In [101]:

```
#grid_search_et=GridSearchCV(estimator=et,param_grid=grid_param_et,cv=3,verbose=1)
```

In [102]:

```
#grid_search_et.fit(X_train,y_train)
```



In [103]:

```
et_tpot_classifier = TPOTClassifier(generations= 5, population_size= 28, offspring_size= 14,
                                   verbosity= 2, early_stop= 12,
                                   config_dict={'sklearn.ensemble.ExtraTreesClassifier': grid
                                   cv = 5, scoring = 'accuracy')
```

In [104]:

```
et_tpot_classifier.fit(X_train,y_train)
```

Generation 1 - Current best internal CV score: 0.80616601653928

Generation 2 - Current best internal CV score: 0.8115597534964779

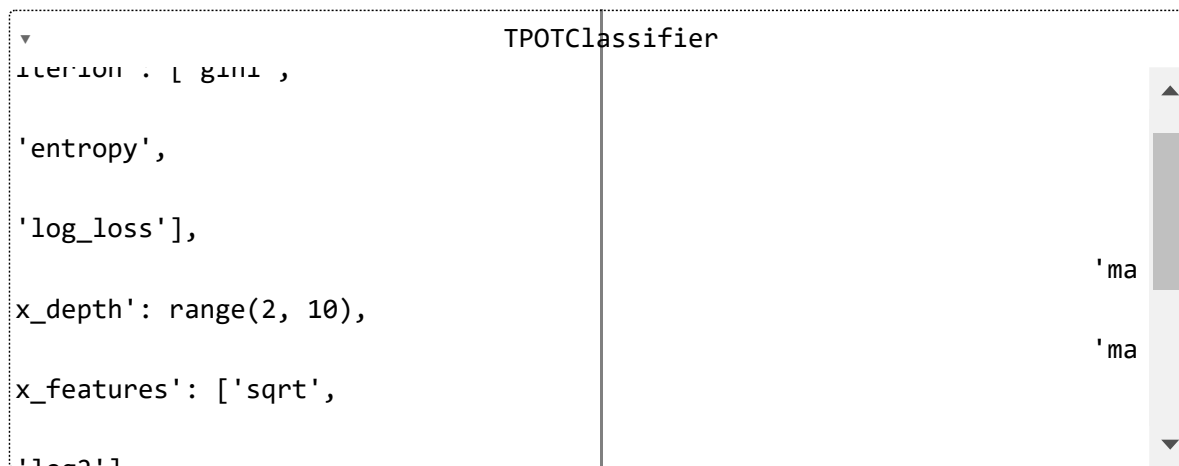
Generation 3 - Current best internal CV score: 0.8115597534964779

Generation 4 - Current best internal CV score: 0.8115597534964779

Generation 5 - Current best internal CV score: 0.8132422114200303

Best pipeline: ExtraTreesClassifier(ExtraTreesClassifier(ExtraTreesClassifier(CombinedDFs(input\_matrix, input\_matrix), criterion=gini, max\_depth=7, max\_features=log2, min\_samples\_leaf=2, min\_samples\_split=6, n\_estimators=100), criterion=gini, max\_depth=6, max\_features=sqrt, min\_samples\_leaf=6, min\_samples\_split=3, n\_estimators=140), criterion=entropy, max\_depth=9, max\_features=sqrt, min\_samples\_leaf=2, min\_samples\_split=5, n\_estimators=140)

Out[104]:



In [105]:

```
et_accuracy = et_tpot_classifier.score(X_test, y_test)
print(et_accuracy)
```

0.8086196503918023

In [106]:

```
et_g_pred=et_tpot_classifier.predict(X_test)
```

In [107]:

```
print(accuracy_score(y_test,et_g_pred))
```

0.8086196503918023

## AdaBoostClassifier

In [108]:

```
from sklearn.ensemble import AdaBoostClassifier
```

In [109]:

```
ad=AdaBoostClassifier()
```

In [110]:

```
ad.fit(X_train,y_train)
```

Out[110]:

```
▼ AdaBoostClassifier  
AdaBoostClassifier()
```

In [111]:

```
ad.score(X_train,y_train)
```

Out[111]:

0.8135886777513855

In [112]:

```
ad_pred=ad.predict(X_test)
```

In [113]:

```
print(accuracy_score(y_test,ad_pred))
```

0.8089210367691381

In [114]:

```
grid_param_ad = {  
    'n_estimators':range(50,100,10),  
    'learning_rate':[0.01,1.0,2.0,5.0,10.0,100.0],  
    'algorithm':['SAMME','SAMME.R']  
}
```

In [115]:


```
grid_search_ad=GridSearchCV(estimator=ad,param_grid=grid_param_ad,cv=3,verbose=1)
```

In [116]:

```
grid_search_ad.fit(X_train,y_train)
```

Fitting 3 folds for each of 60 candidates, totalling 180 fits

Out[116]:



In [117]:

```
grid_search_ad.best_params_
```

Out[117]:

```
{'algorithm': 'SAMME.R', 'learning_rate': 1.0, 'n_estimators': 90}
```


In [118]:

```
grid_search_ad_with_best_parm = AdaBoostClassifier(n_estimators =90,algorithm='SAMME.R', le
```

In [119]:

```
grid_search_ad_with_best_parm.fit(X_train,y_train)
```

Out[119]:



In [120]:

```
grid_search_ad_with_best_parm_pred = grid_search_ad_with_best_parm.predict(X_test)
```

In [121]:

```
print(accuracy_score(y_test,grid_search_ad_with_best_parm_pred))
```

```
0.8081173397629093
```

## VotingClassifier

In [122]:

```
#vt=VotingClassifier(estimators=[('DT',dc), ('RF',rf), ('SVM', sv),('ADABOOST',ad)],voting=
```

In [125]:

```
#vt.fit(X_train,y_train)
from sklearn.linear_model import LogisticRegression
```

In [126]:

```
estimaors=[
    ('RFC',RandomForestClassifier()),
    ('Lr',LogisticRegression()),
    ('ETC',ExtraTreesClassifier(criterion='gini',max_depth=10,min_samples_leaf=2,min_sample
    ('BGC',BaggingClassifier(base_estimator=LogisticRegression())))
]
```

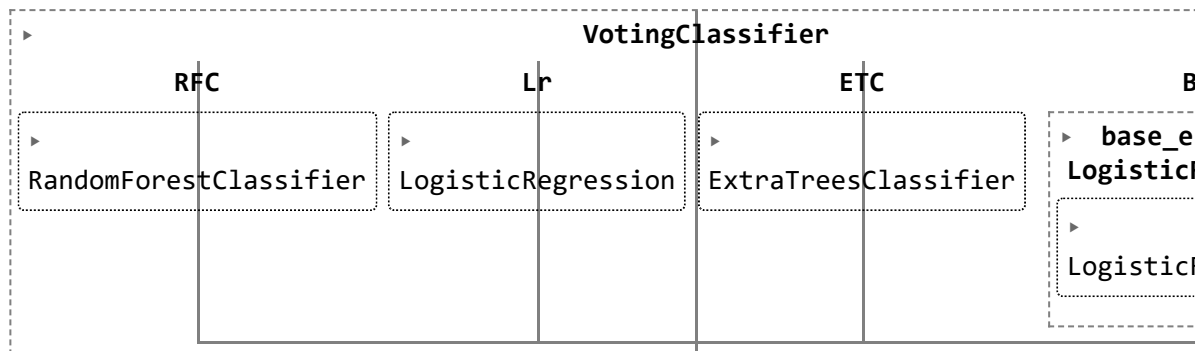
In [131]:

```
vt = VotingClassifier(estimators=estimaors,voting='hard')
```

In [132]:

```
vt.fit(X_train,y_train)
```

Out[132]:



In [133]:

```
vt.score(X_train,y_train)
```

Out[133]:

```
0.7519794140934284
```

In [134]:

```
vt_pred = vt.predict(X_test)
```

In [135]:

```
print(accuracy_score(y_test,vt_pred))
```

```
0.7492465340566606
```

In [ ]:

